

OPENSIFT CONTAINER PLATFORM

TECHNICAL OVERVIEW



linkedin.com/company/red-hat



facebook.com/redhatinc



youtube.com/user/RedHatVideos



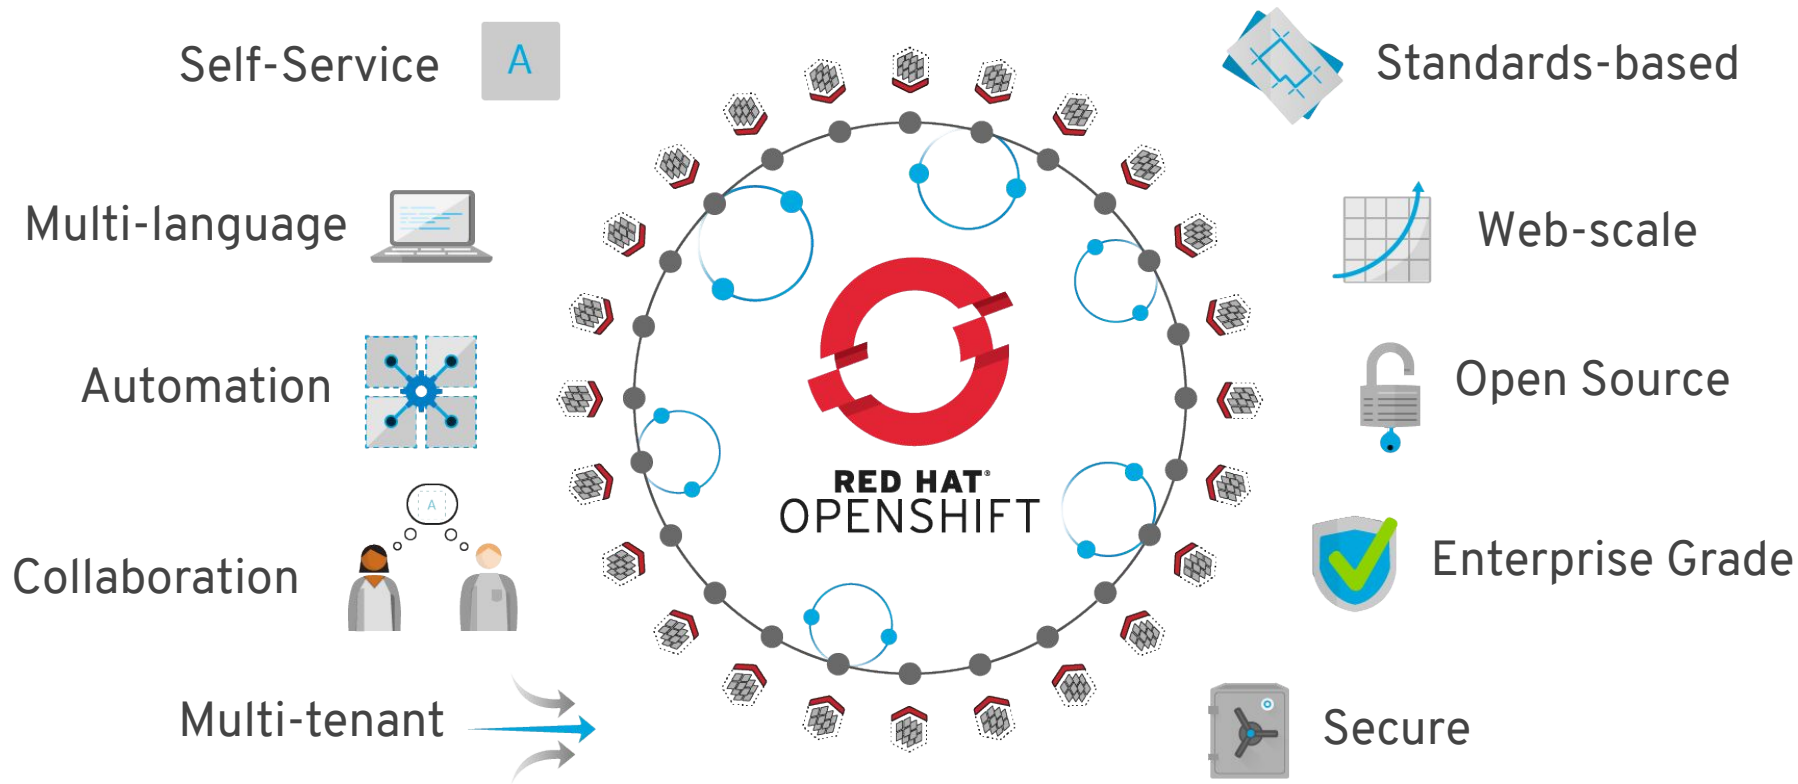
twitter.com/RedHat

PRESENTER NAME
PRESENTER TITLE
DATE

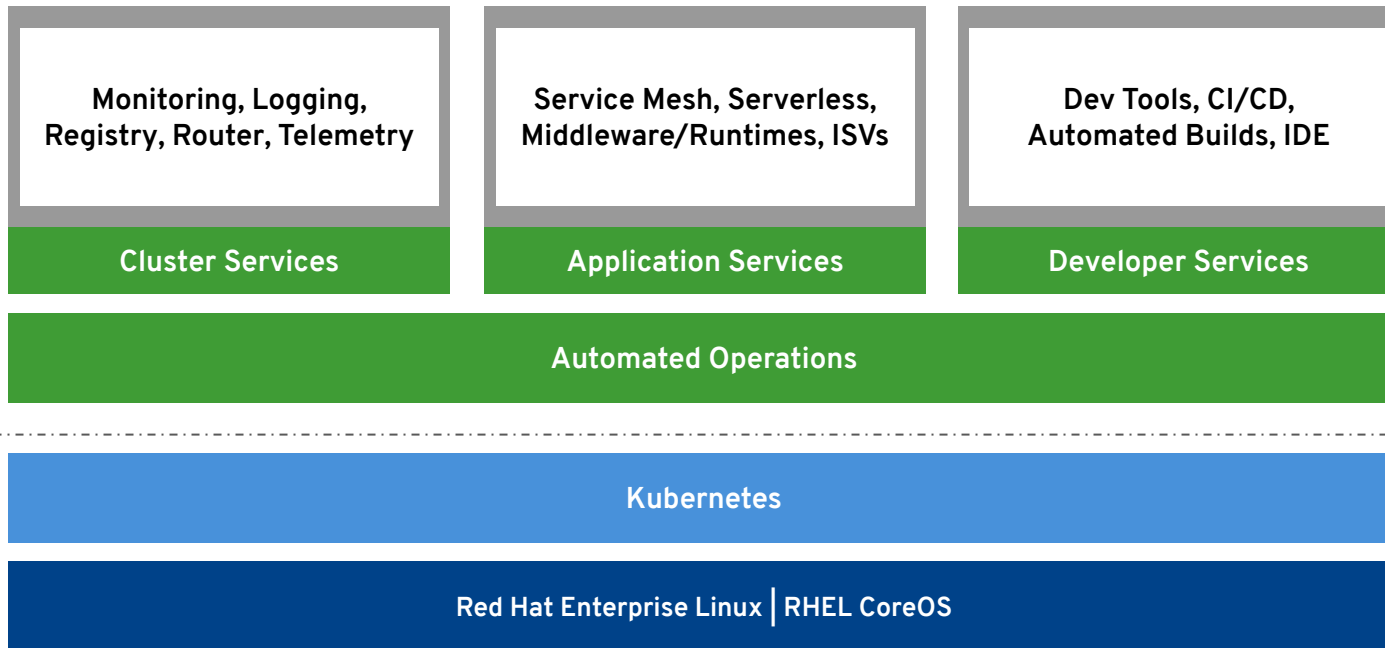




Functional overview



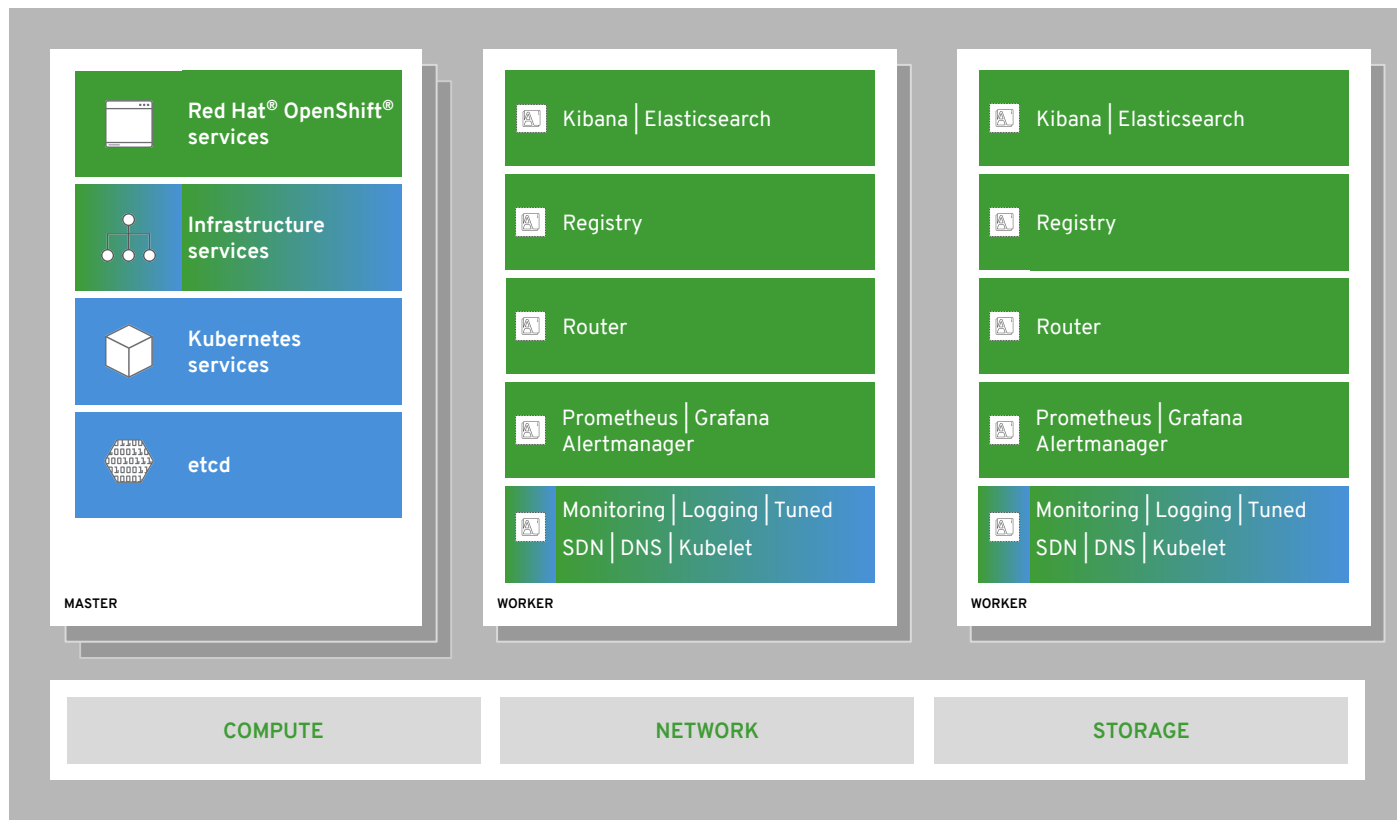
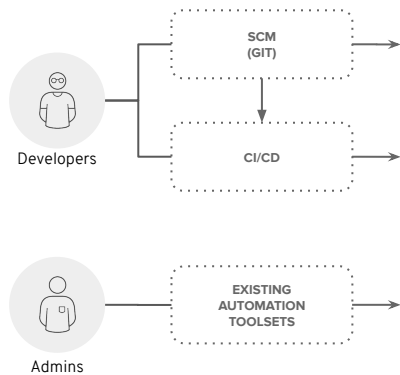
Value of OpenShift



Best IT Ops Experience

CaaS ↔ PaaS ↔ FaaS

Best Developer Experience



Database

Streaming & Messaging

Application Definition & Image Build

Continuous Integration & Delivery

Platform

Observability and Analysis

App Definition and Development



Orchestration & Management



Cloud-Native Storage

Container Runtime

Cloud-Native Network

Runtime



Automation & Configuration

Container Registry

Security & Compliance

Key Management

Provisioning




Public

Kubernetes Certified Service Provider

Kubernetes Training Partner





OpenShift and Kubernetes core concepts

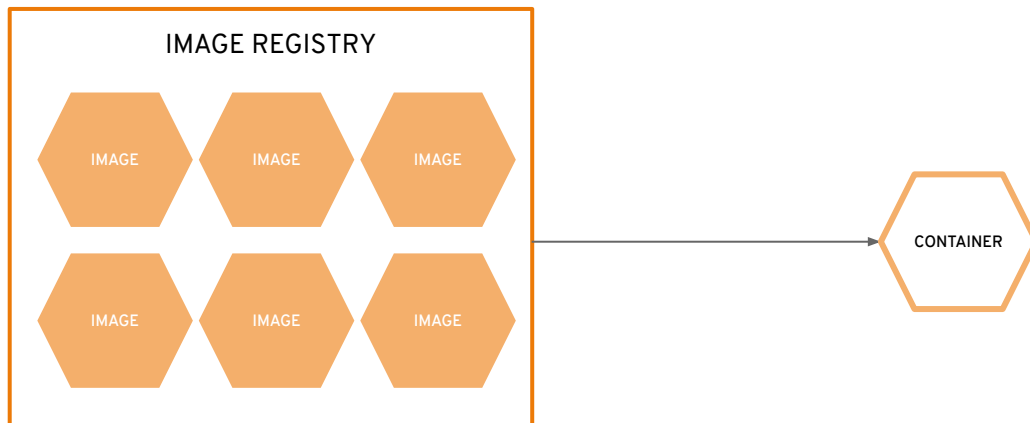
a container is the smallest compute unit



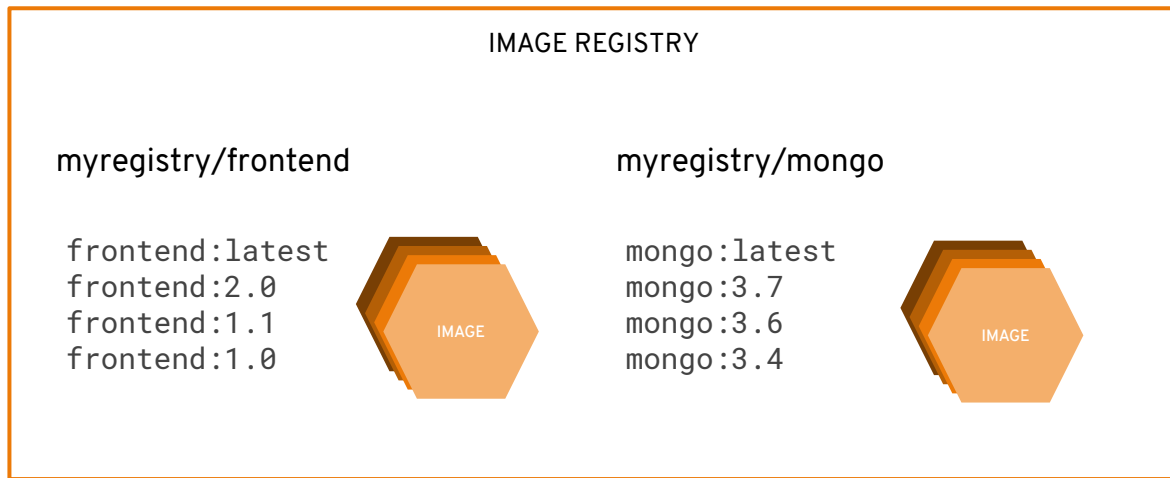
containers are created from container images



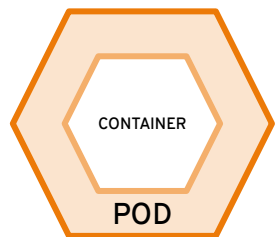
container images are stored in an image registry



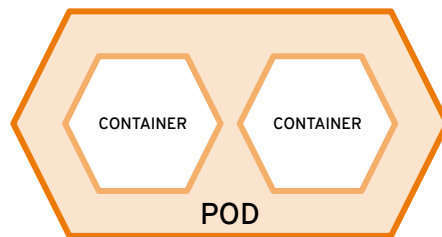
an image repository contains all versions of an image in the image registry



containers are wrapped in pods which are units of deployment and management

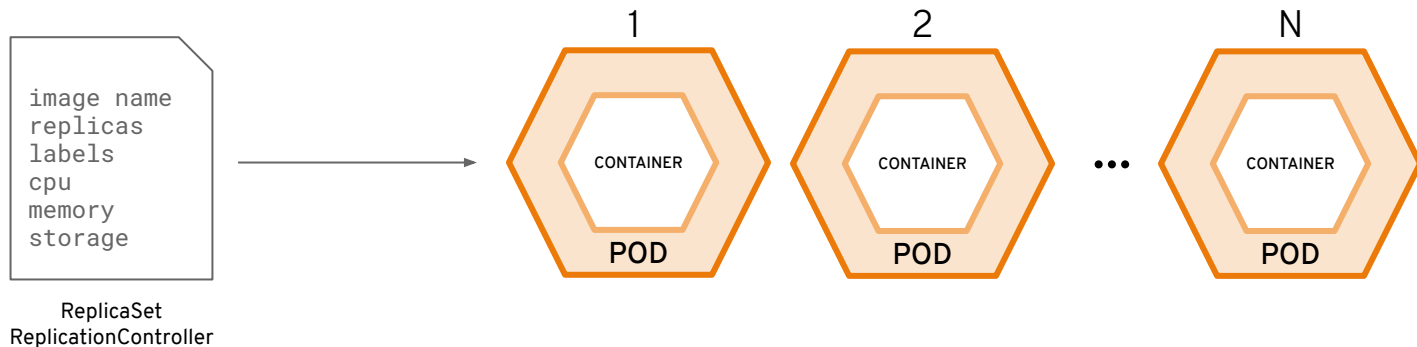


10.140.4.44

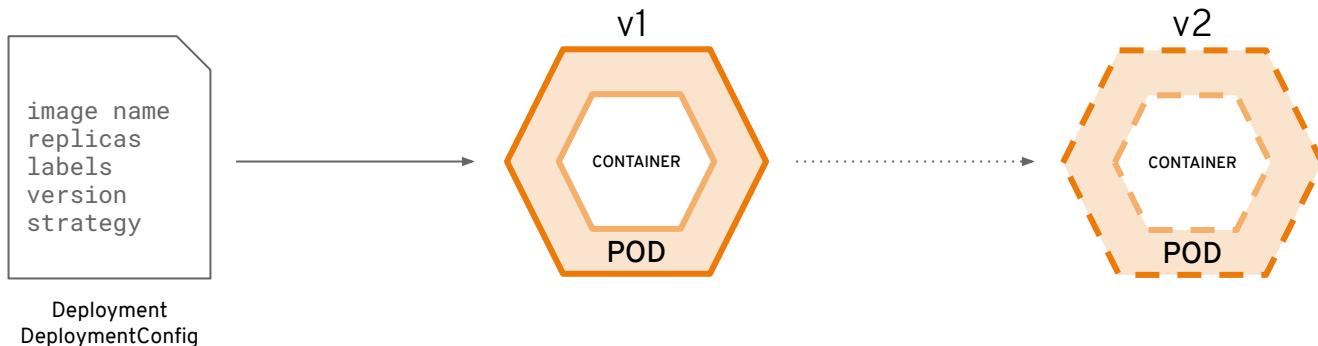


10.15.6.55

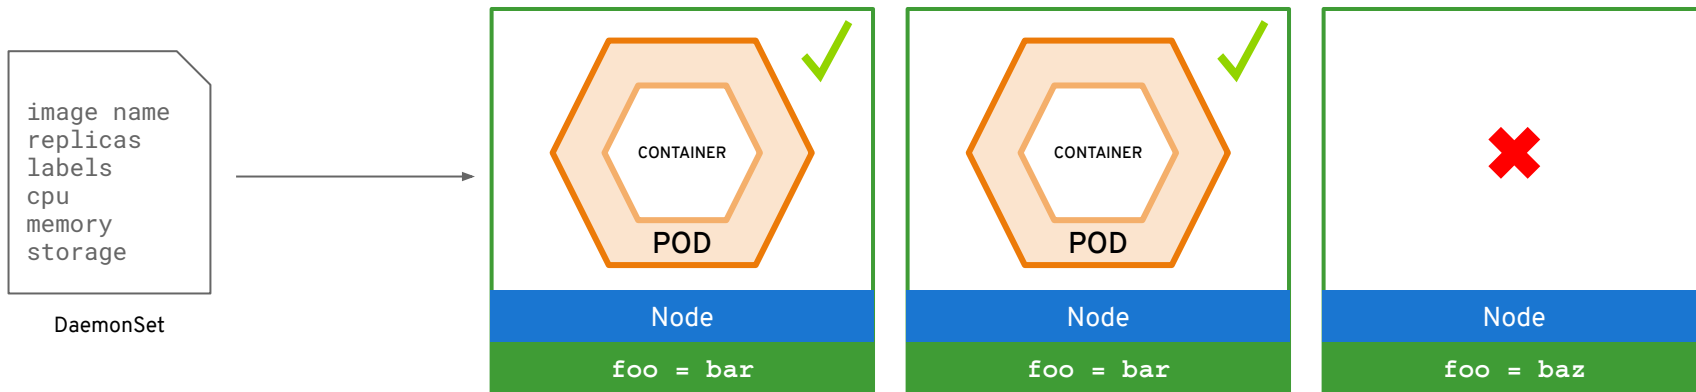
ReplicationControllers & ReplicaSets ensure a specified number of pods are running at any given time



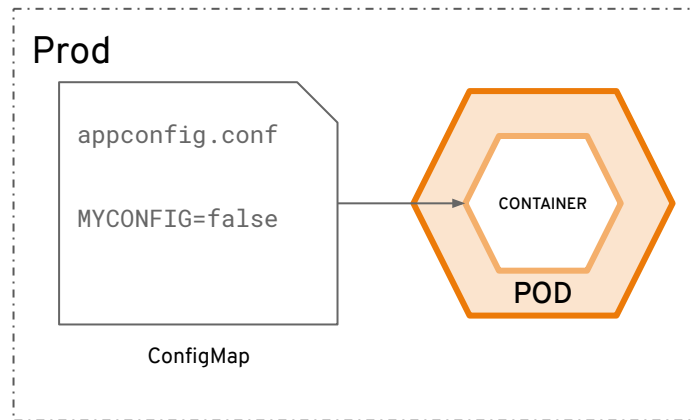
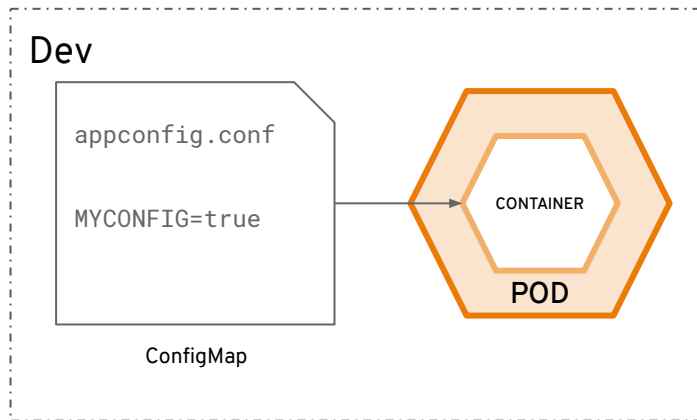
Deployments and DeploymentConfigurations define how to roll out new versions of Pods



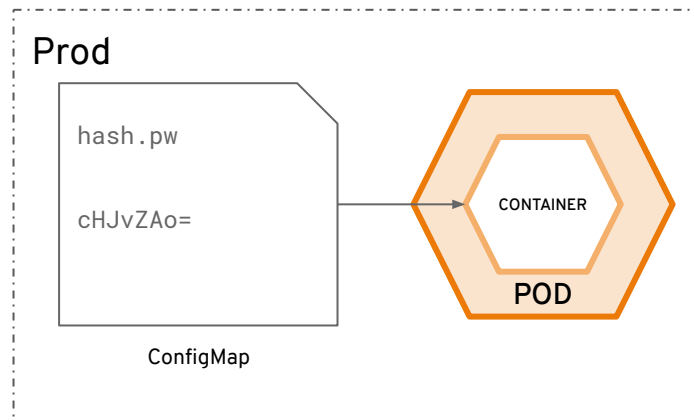
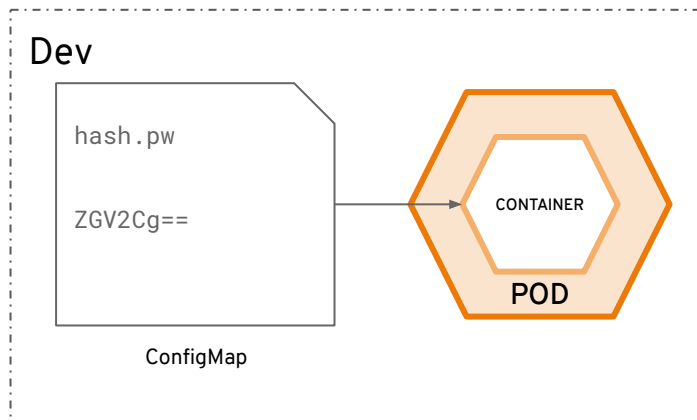
a daemonset ensures that all
(or some) nodes run a copy of a pod



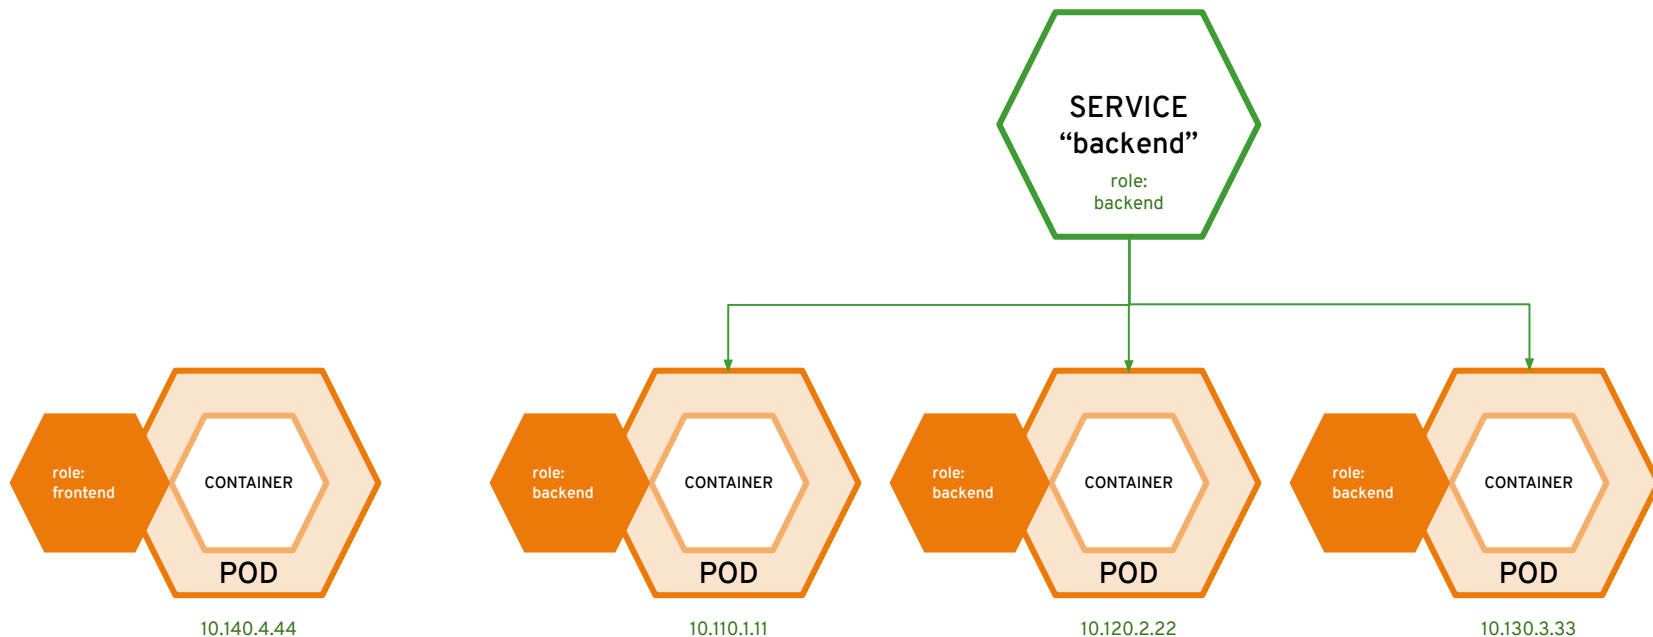
configmaps allow you to decouple configuration artifacts from image content



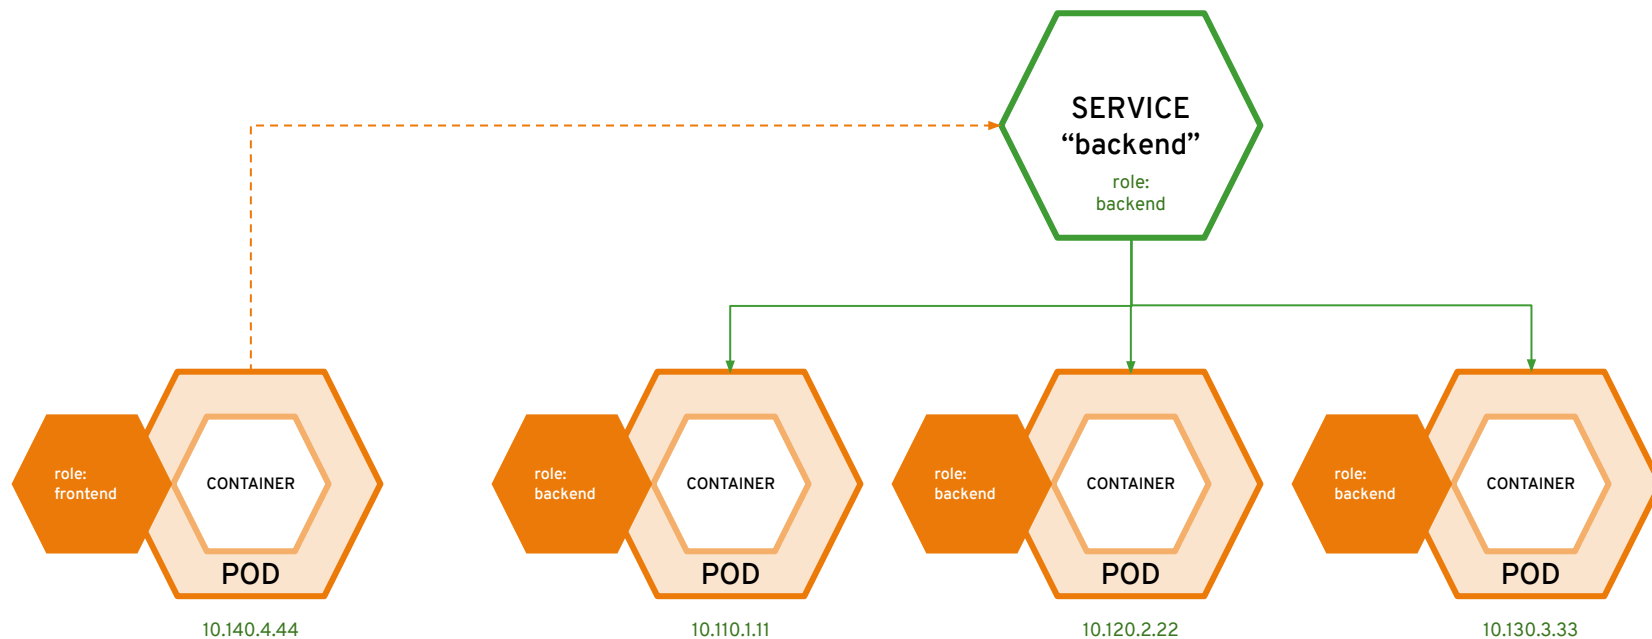
secrets provide a mechanism to hold sensitive information such as passwords



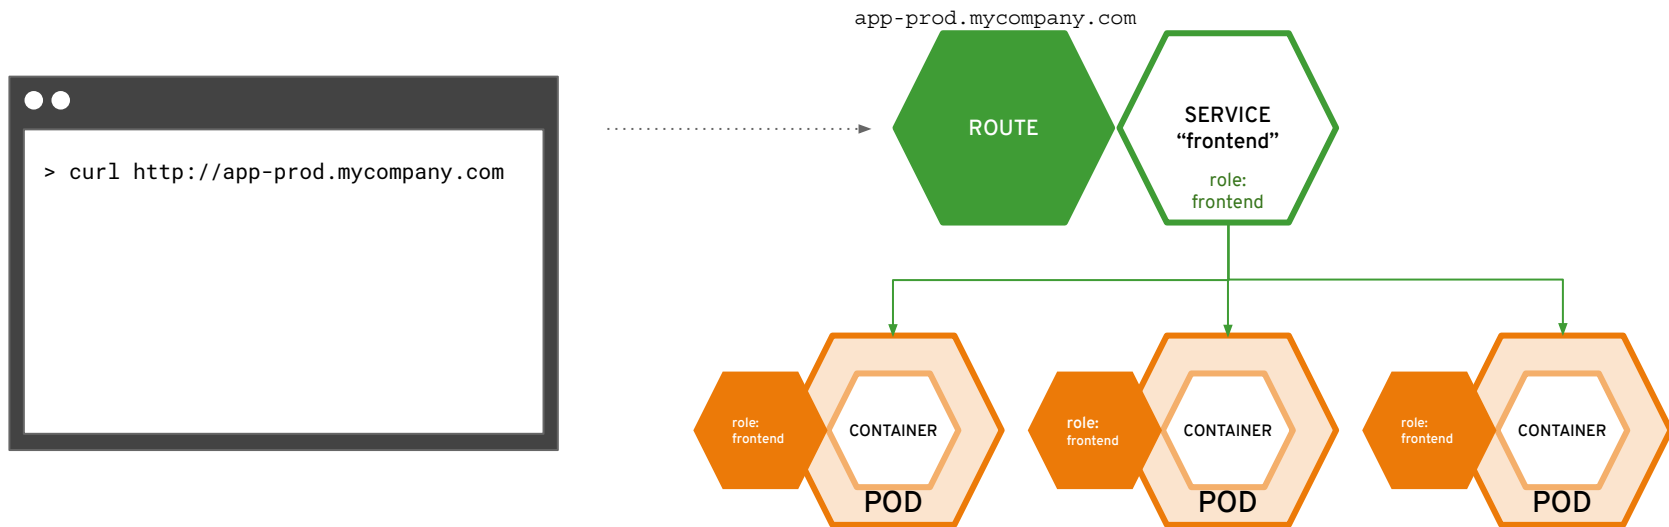
services provide internal load-balancing and service discovery across pods



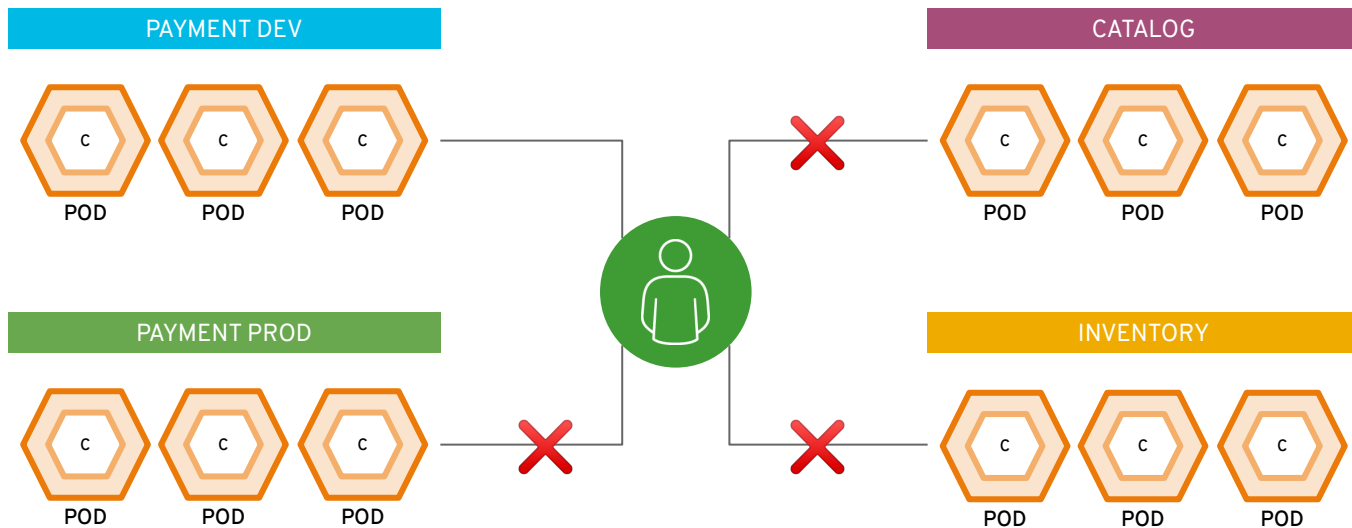
apps can talk to each other via services



routes make services accessible to clients outside the environment via real-world urls



projects isolate apps across environments, teams, groups and departments





OpenShift 4 Architecture

your choice of infrastructure

COMPUTE

NETWORK

STORAGE

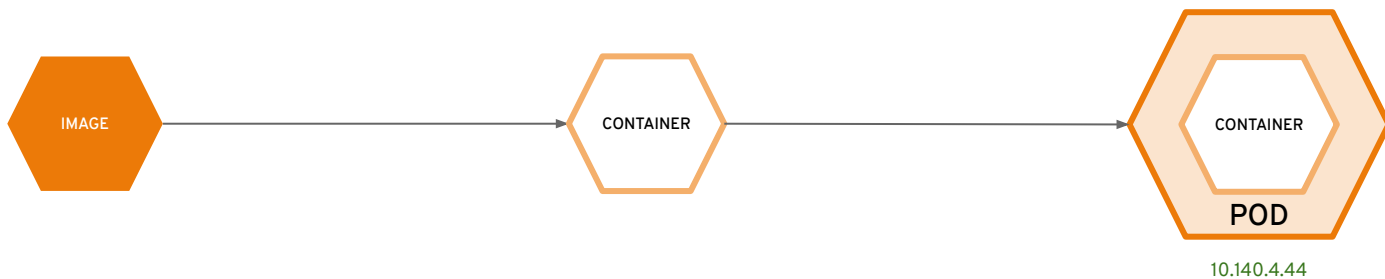
workers run workloads



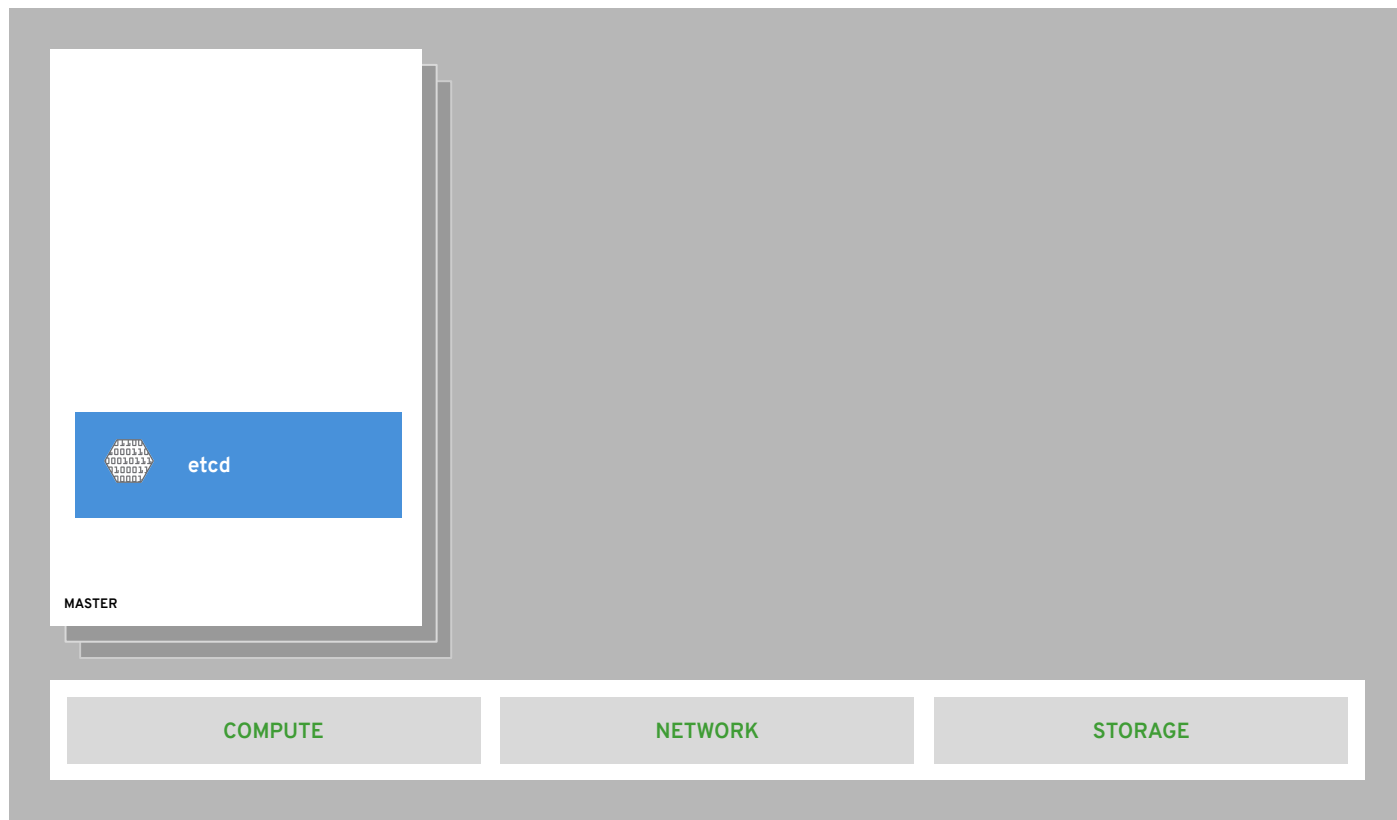
masters are the control plane



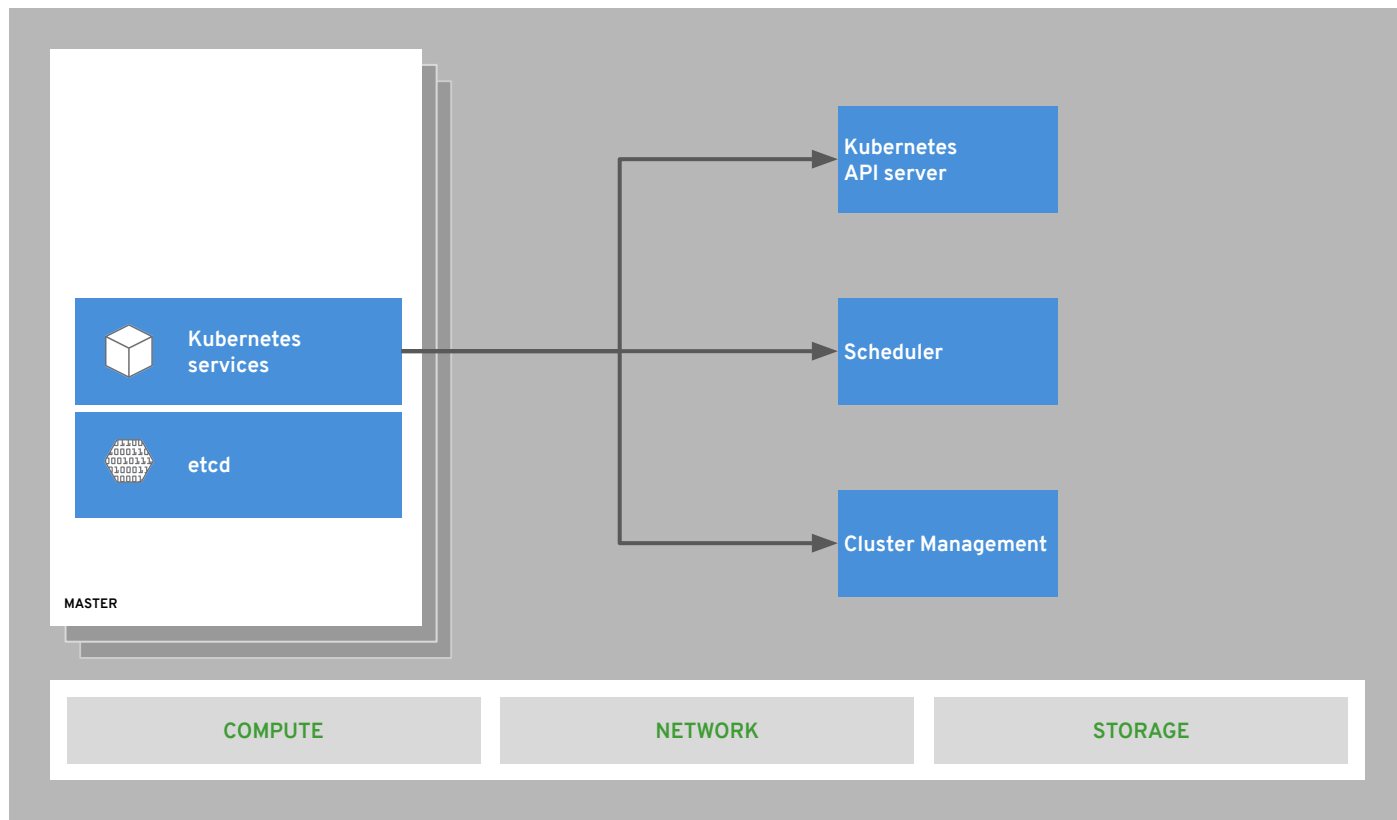
everything runs in pods



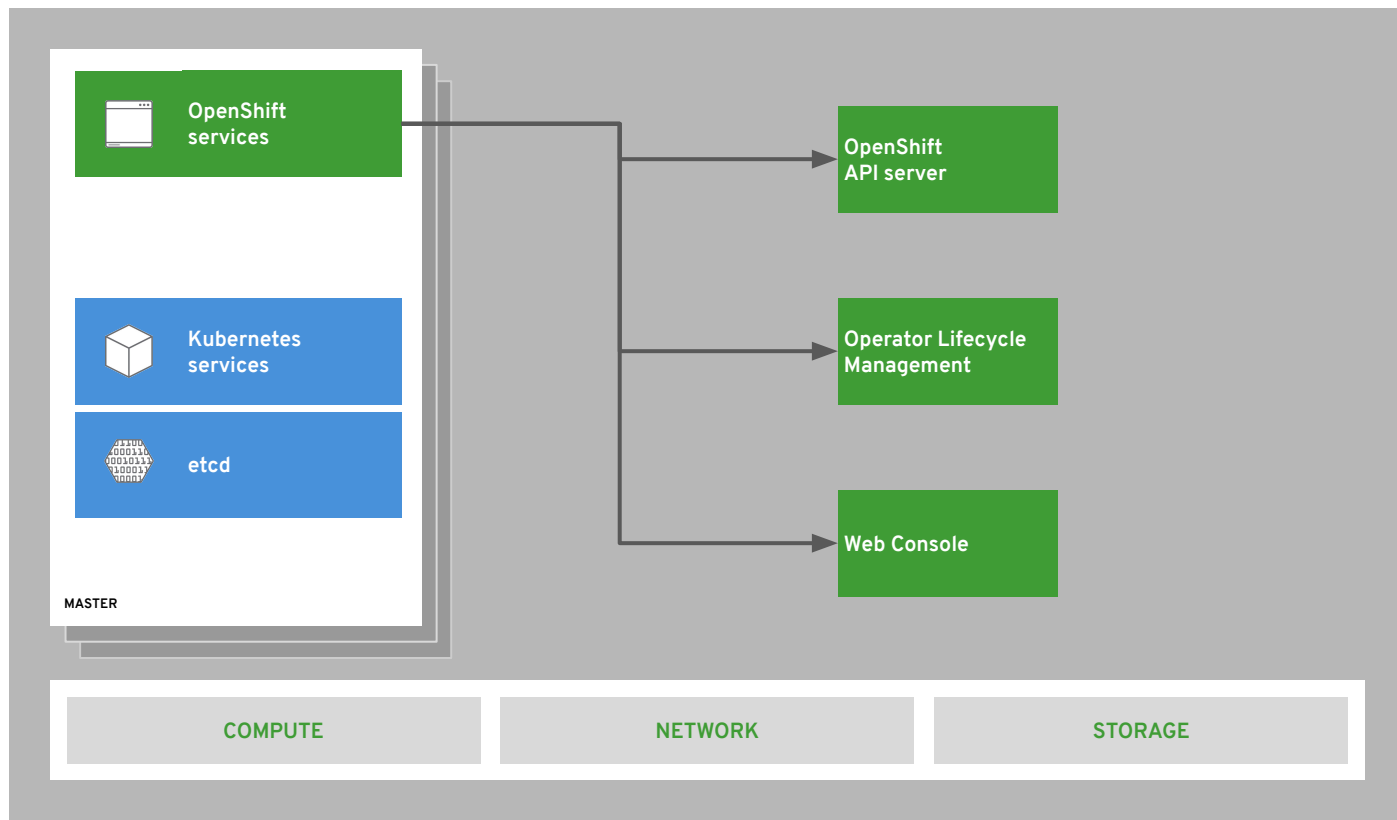
state of everything



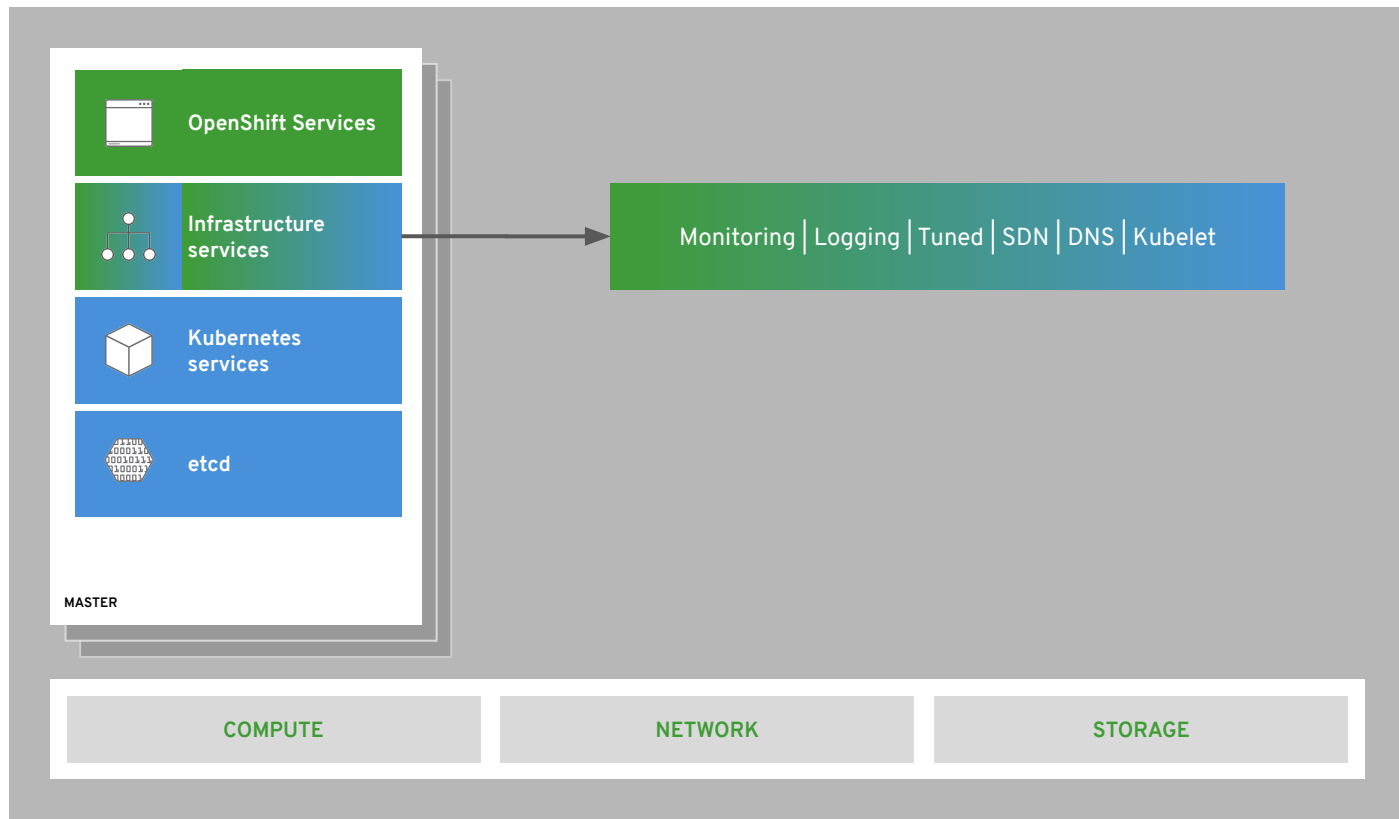
core kubernetes components



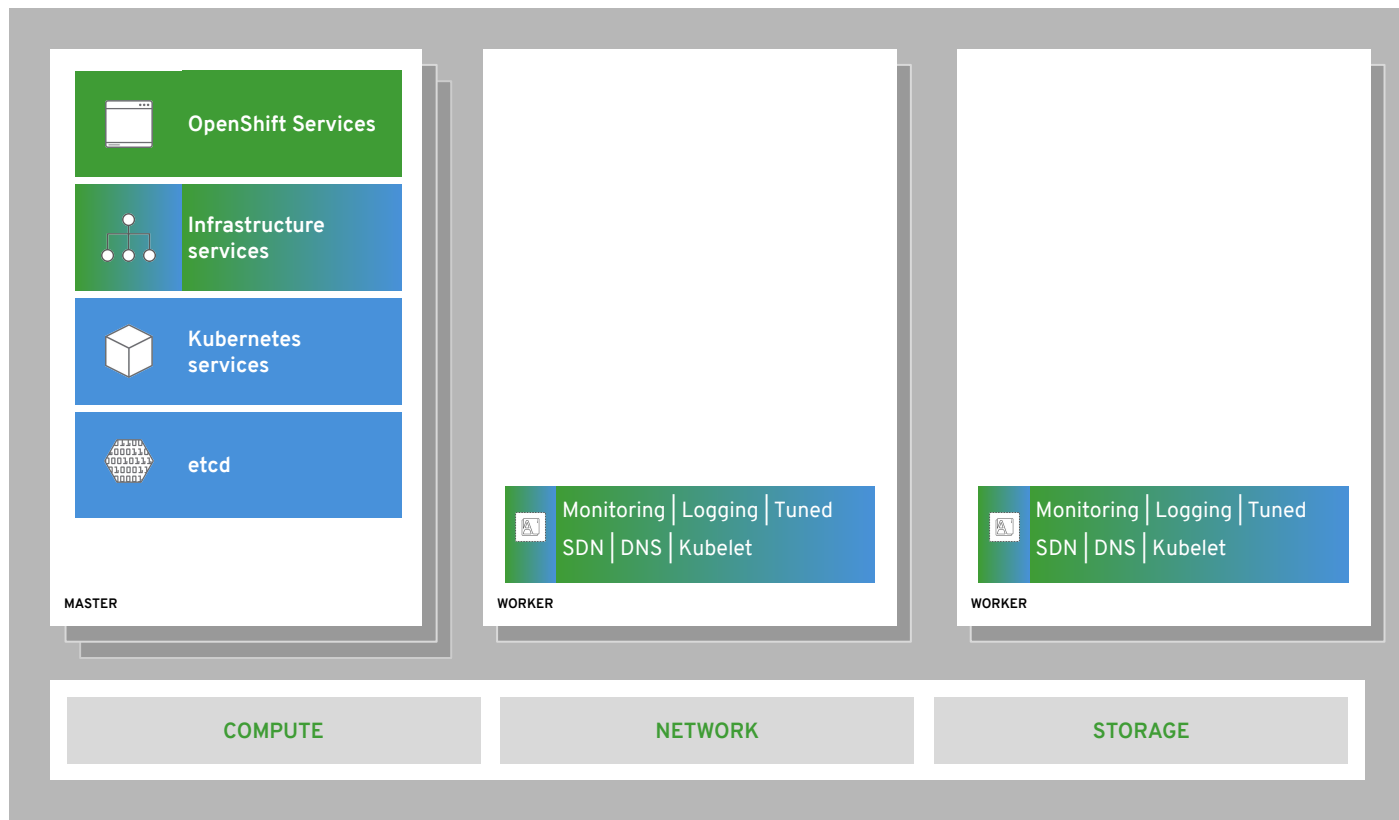
core OpenShift components



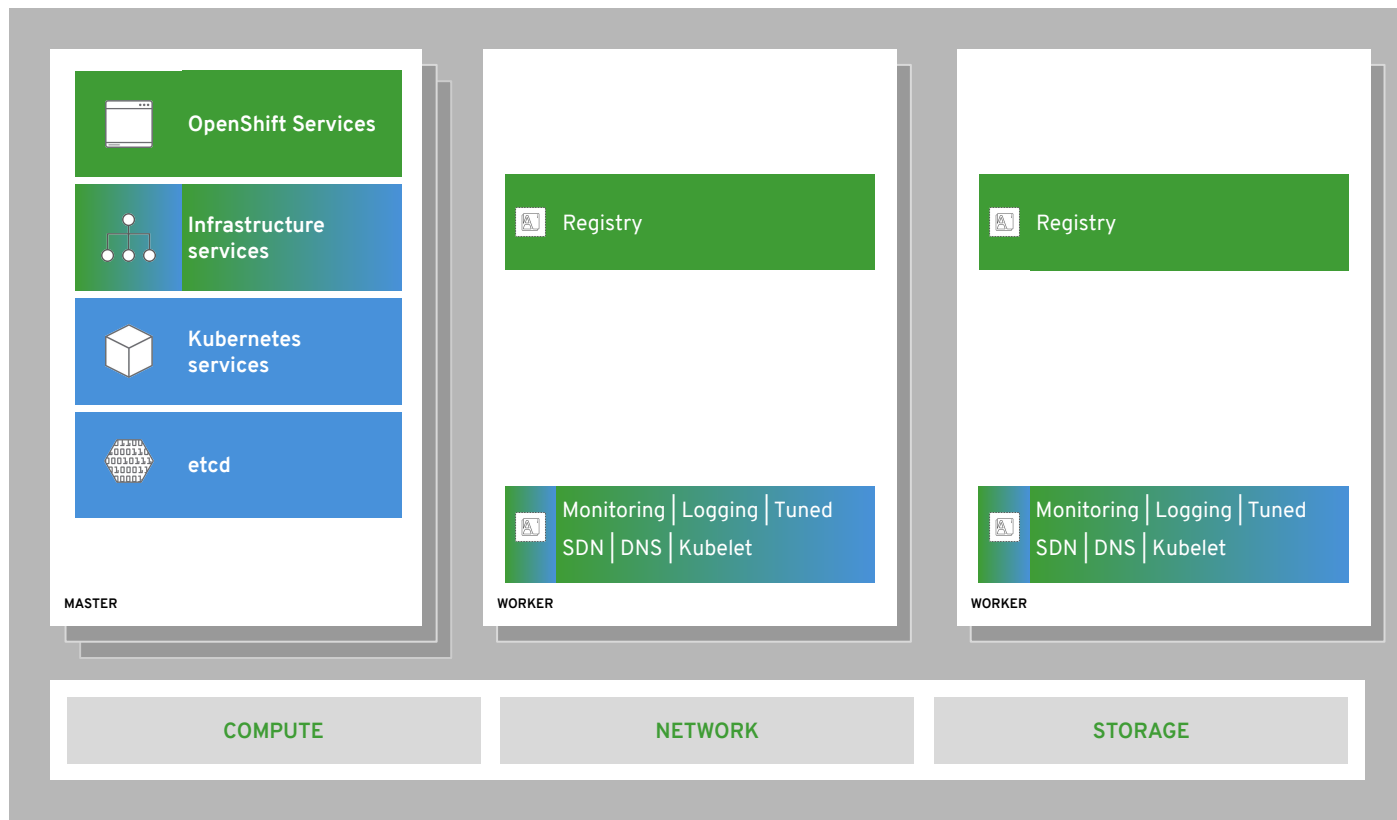
internal and support infrastructure services



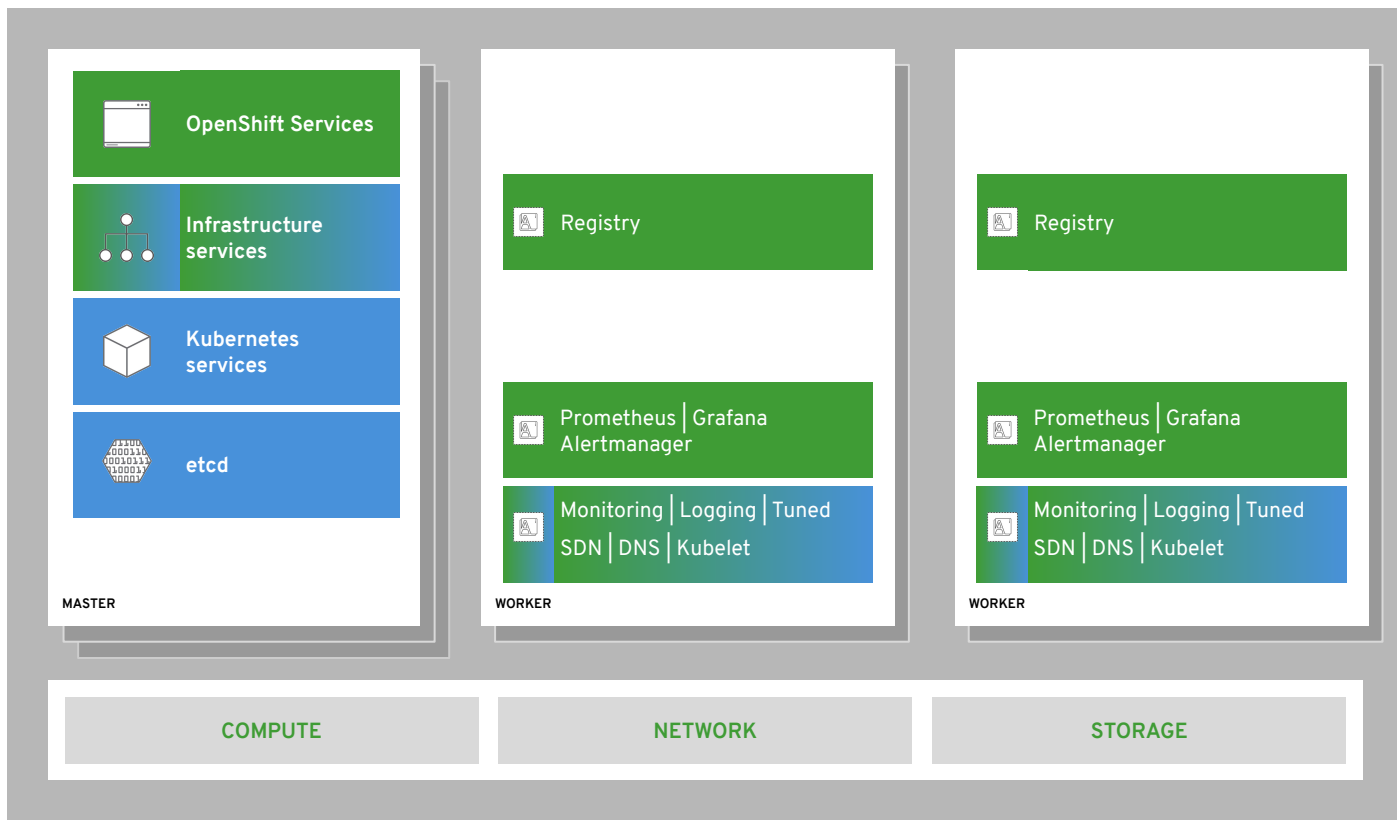
run on all hosts



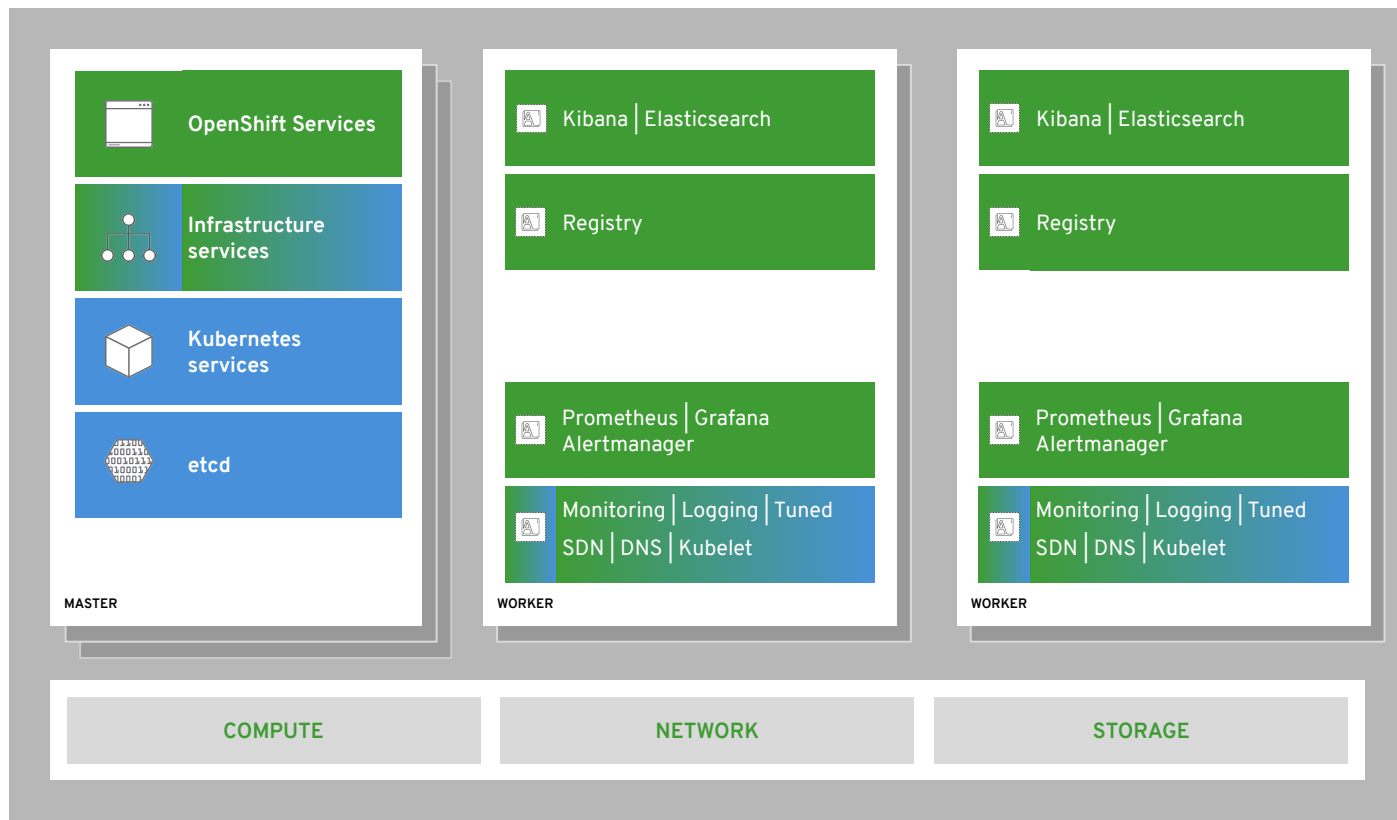
integrated image registry



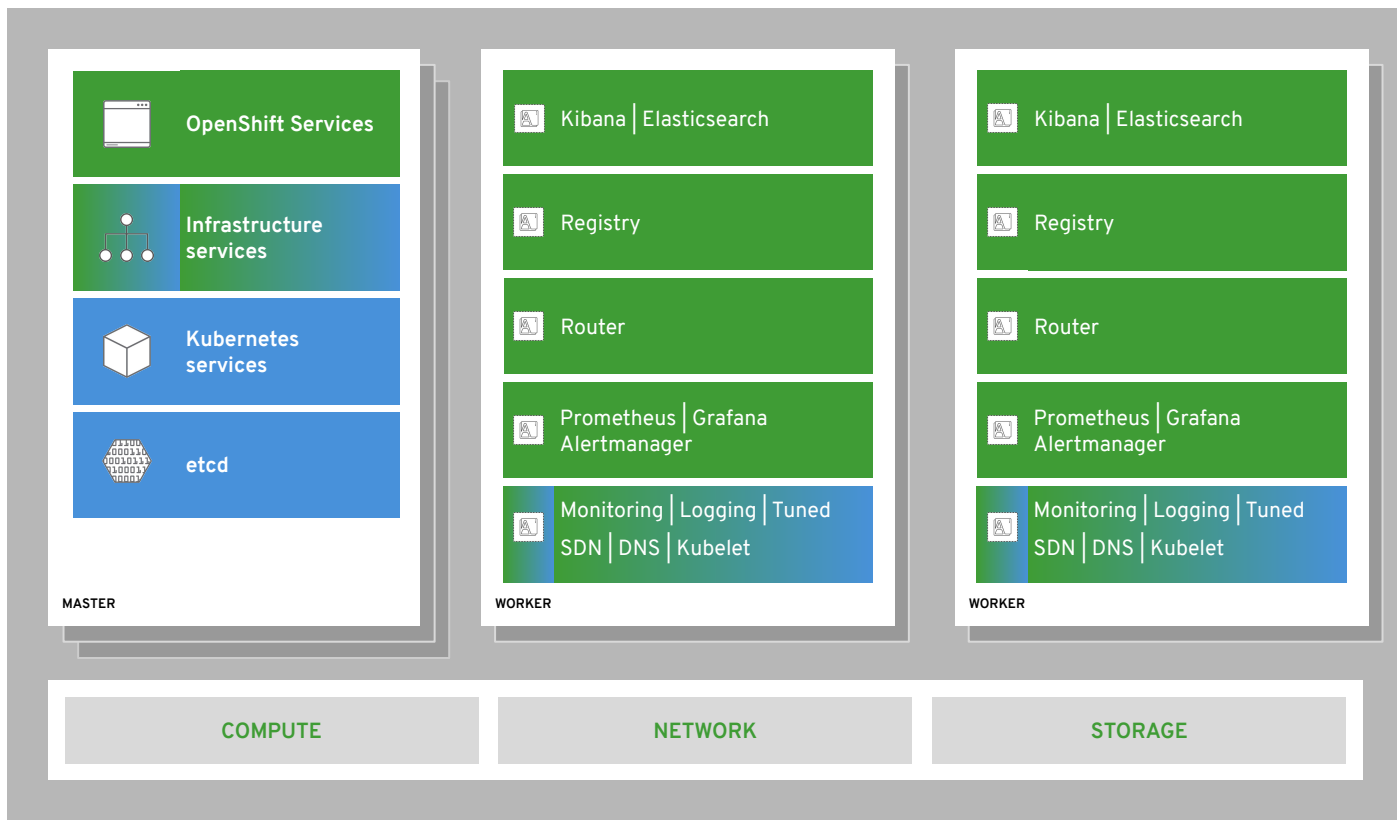
cluster monitoring



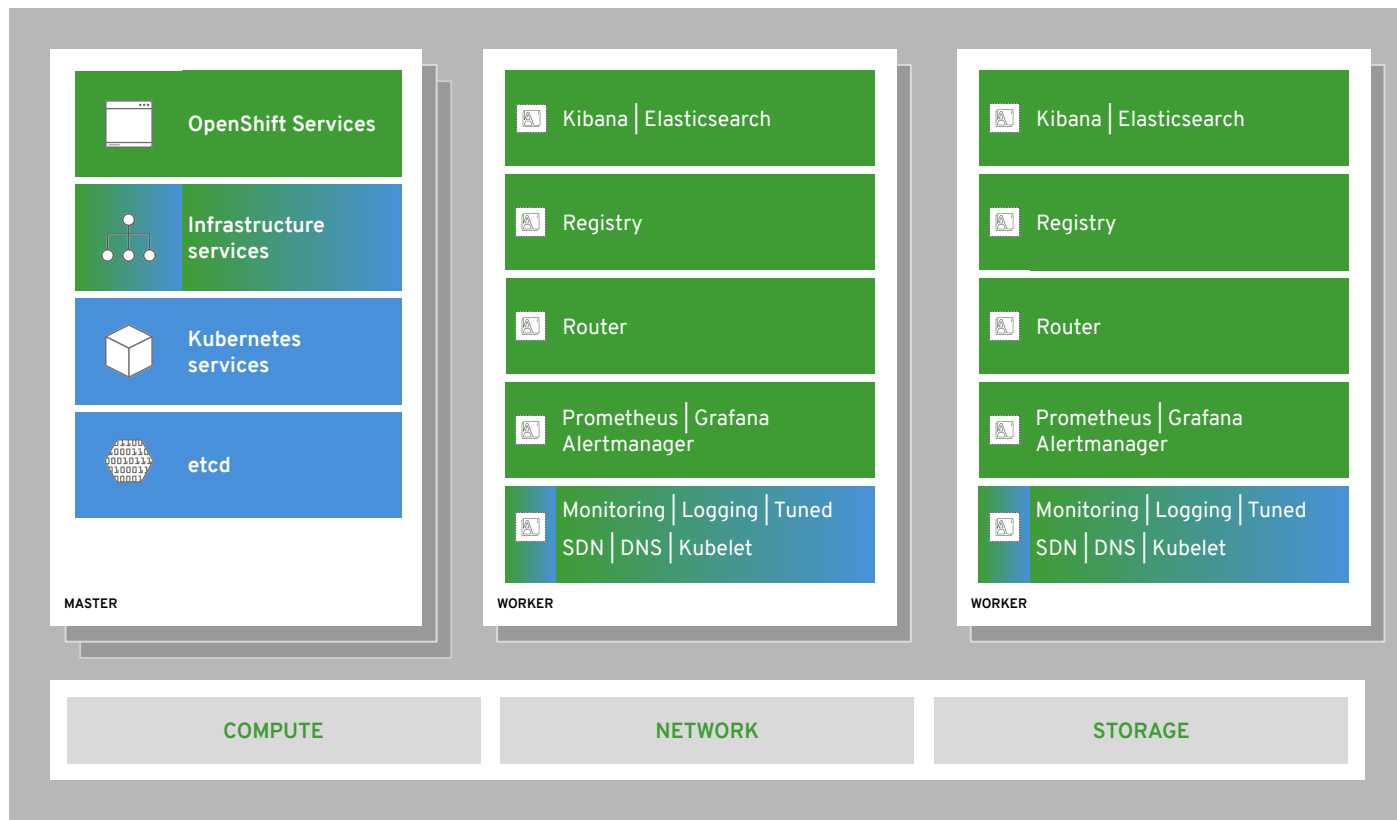
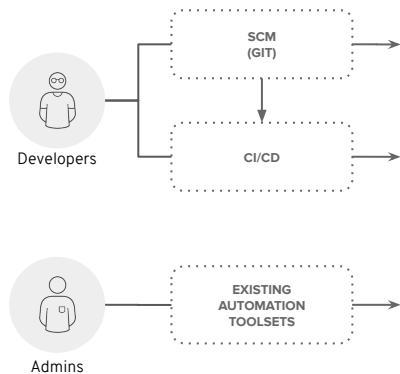
log aggregation



integrated routing



dev and ops via web, cli, API, and IDE





OpenShift lifecycle, installation & upgrades

OpenShift 4 Installation

Two new paradigms for
deploying clusters

Installation Paradigms

OPENSIFT CONTAINER PLATFORM

Full Stack Automated

Simplified opinionated “Best Practices” for cluster provisioning

Fully automated installation and updates including host container OS.



Pre-existing Infrastructure

Customer managed resources & infrastructure provisioning

Plug into existing DNS and security boundaries



HOSTED OPENSIFT

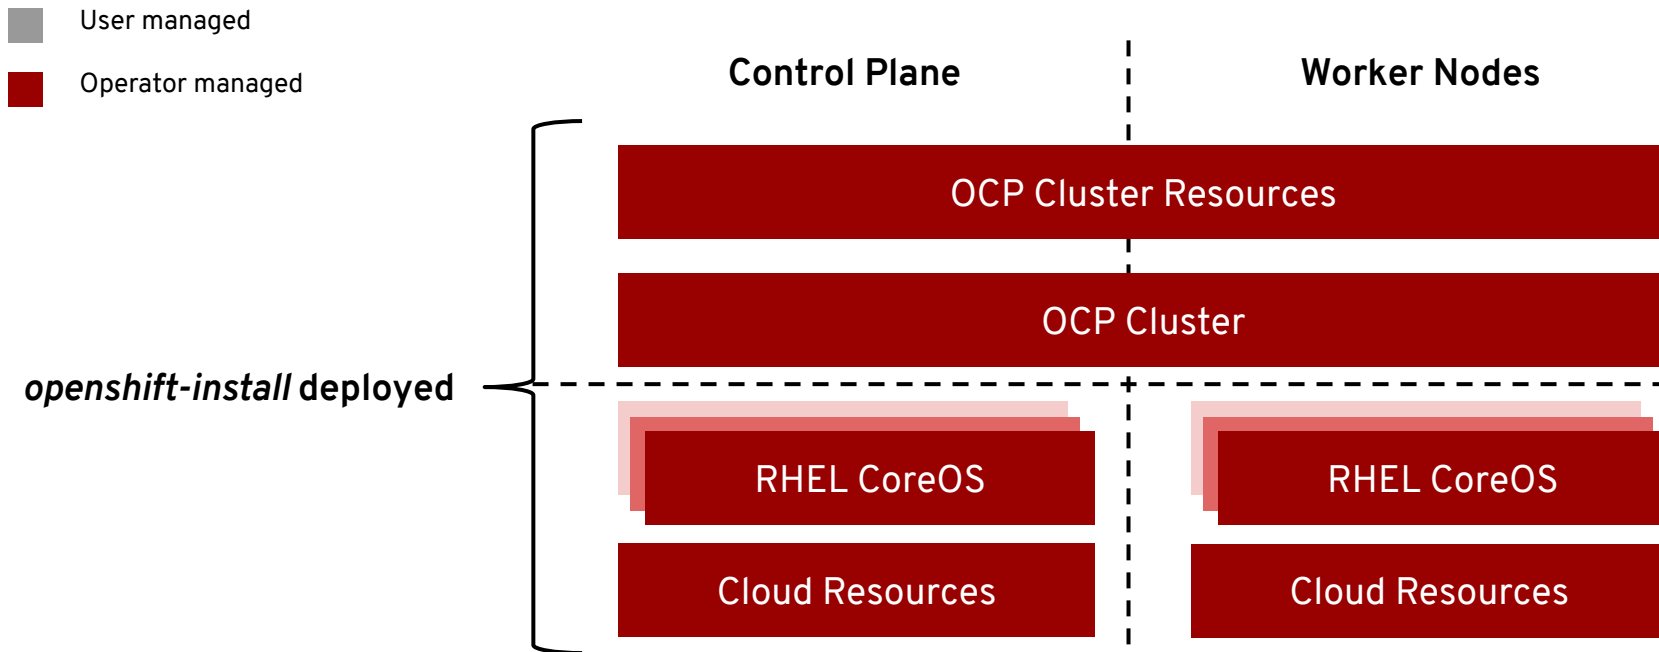
Azure Red Hat OpenShift

Deploy directly from the Azure console. Jointly managed by Red Hat and Microsoft Azure engineers.

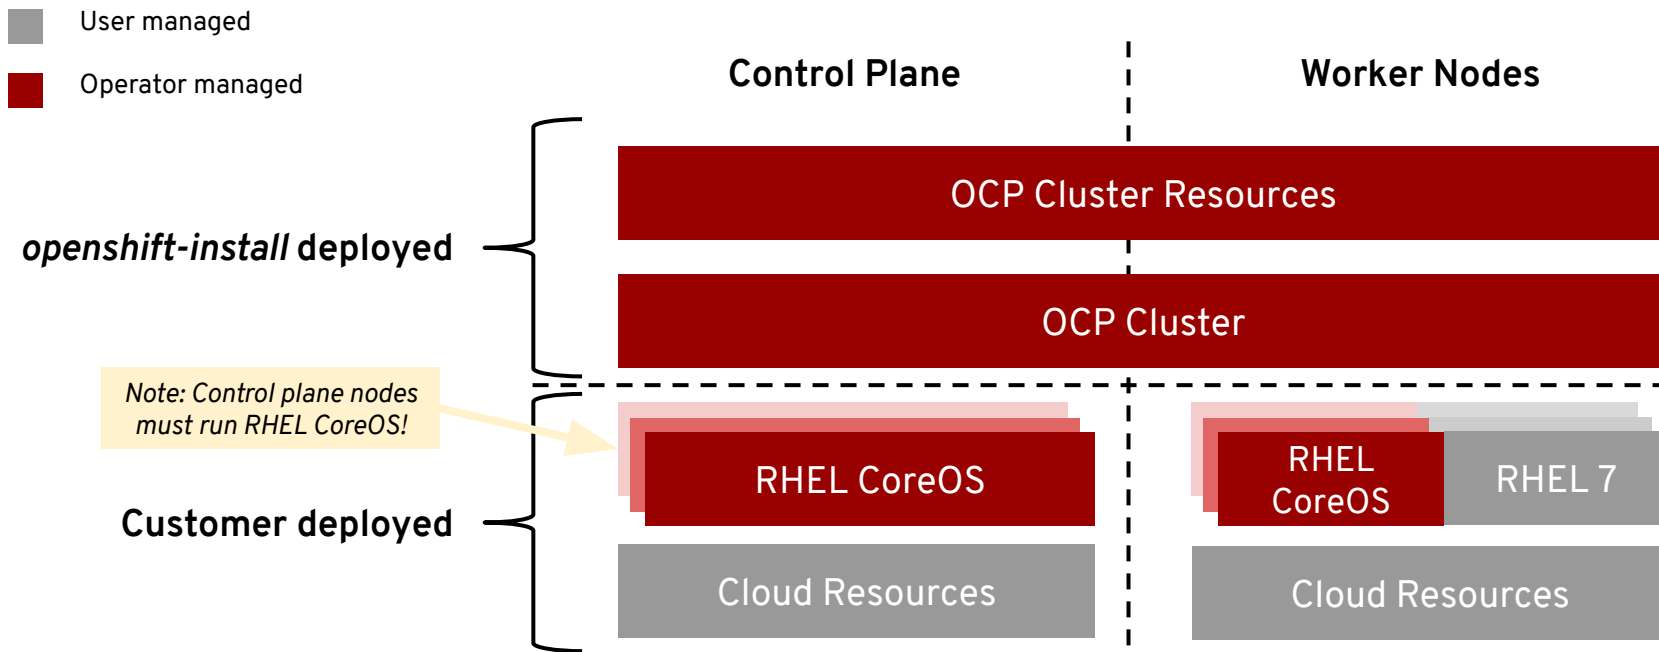
OpenShift Dedicated

Get a powerful cluster, fully Managed by Red Hat engineers and support.

Full-stack Automated Installation



Pre-existing Infrastructure Installation



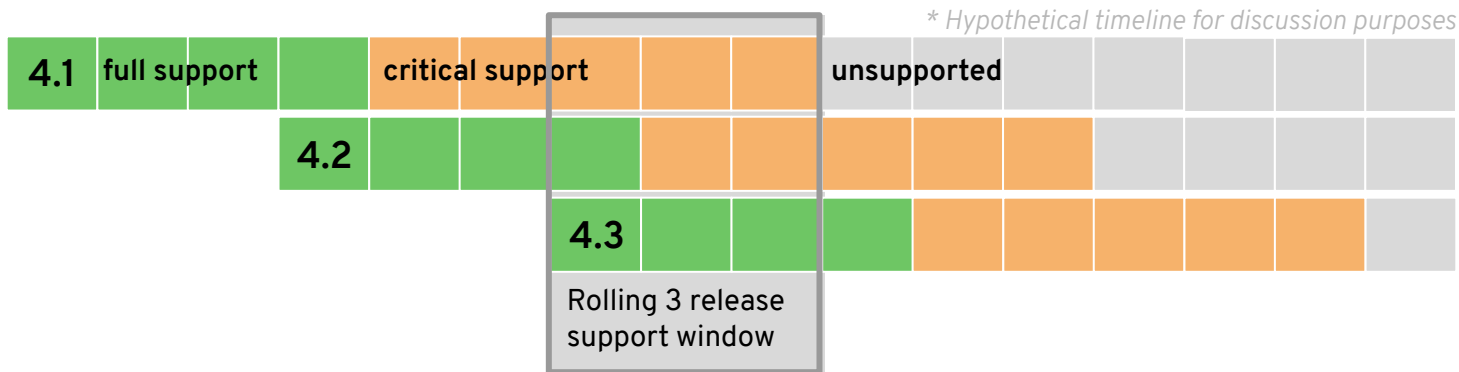
Comparison of Paradigms

	Full Stack Automation	Pre-existing Infrastructure
Build Network	Installer	User
Setup Load Balancers	Installer	User
Configure DNS	Installer	User
Hardware/VM Provisioning	Installer	User
OS Installation	Installer	User
Generate Ignition Configs	Installer	Installer
OS Support	Installer: RHEL CoreOS	User: RHEL CoreOS + RHEL 7
Node Provisioning / Autoscaling	Yes	Only for providers with OpenShift Machine API support

OpenShift 4 Lifecycle

Supported paths for
upgrades and migrations

Support Timelines



New model

Release based, not date based. Rolling three release window for support.

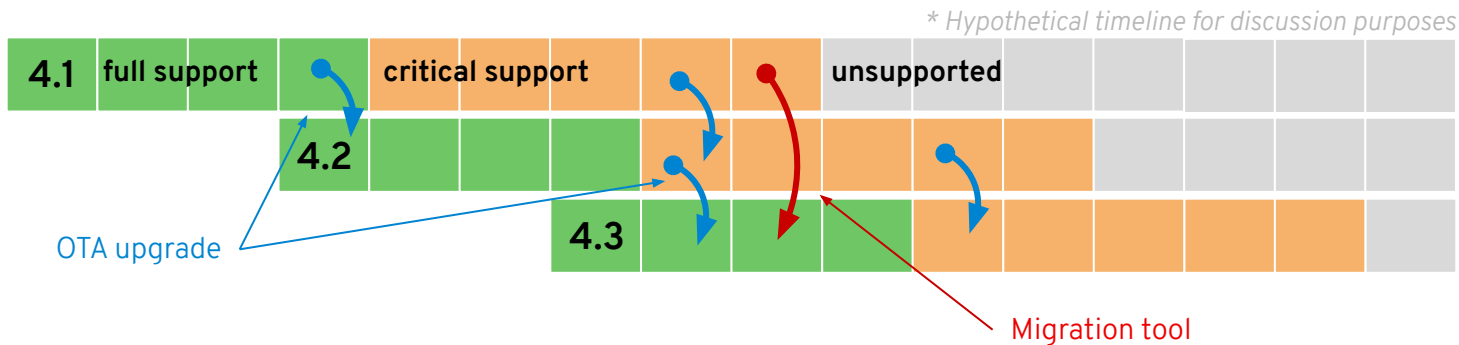
The overall 4 series will be supported for at least three years

- Minimum two years full support (likely more)
- One year maintenance past the end of full support

EUS release planned

Supported for 14 months of critical bug and critical security fixes instead of the normal 5 months. If you stay on the EUS for its entire life, you must use the application migration tooling to move to a new cluster

Upgrades vs. Migrations



OTA Upgrades

Works between two minor releases in a serial manner.

Happy path = migrate through each version

On a regular cadence, migrate to the next supported version.

Optional path = migration tooling

If you fall more than two releases behind, you must use the application migration tooling to move to a new cluster.

Current minor release

Full support for all bugs and security issues
1 month full support overlap with next release to aid migrations

Previous minor release

Fixes for critical bugs and security issues for 5 months



Operations and infrastructure deep dive

Red Hat Enterprise Linux CoreOS

The OpenShift operating
system

Red Hat Enterprise Linux

RED HAT® ENTERPRISE LINUX®

General Purpose OS

RED HAT® ENTERPRISE LINUX CoreOS

Immutable container host

BENEFITS

- 10+ year enterprise life cycle
- Industry standard security
- High performance on any infrastructure
- Customizable and compatible with wide ecosystem of partner solutions

- Self-managing, over-the-air updates
- Immutable and tightly integrated with OpenShift
- Host isolation is enforced via Containers
- Optimized performance on popular infrastructure

WHEN TO USE

When customization and integration with additional solutions is required

When cloud-native, hands-free operations are a top priority

Immutable Operating System

Red Hat Enterprise Linux CoreOS is versioned with OpenShift

CoreOS is tested and shipped in conjunction with the platform. Red Hat runs thousands of tests against these configurations.

Red Hat Enterprise Linux CoreOS is managed by the cluster

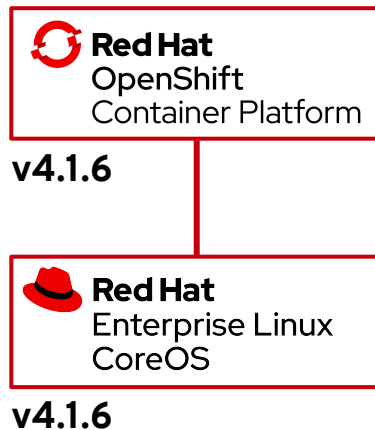
The Operating system is operated as part of the cluster, with the config for components managed by Machine Config

Operator:

- CRI-O config
- Kubelet config
- Authorized registries
- SSH config

RHEL CoreOS admins are responsible for:

Nothing. 😊 🙌





A lightweight, OCI-compliant container runtime

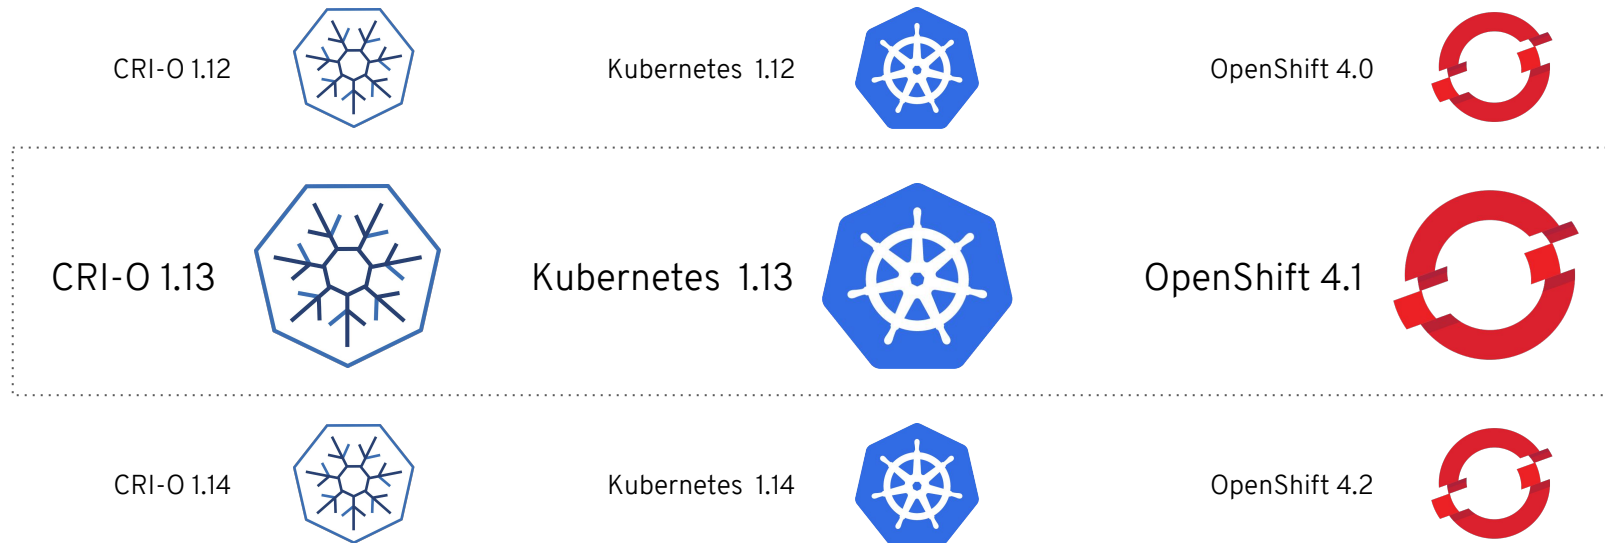
Minimal and Secure
Architecture

Optimized for
Kubernetes

Runs any
OCI-compliant image
(including docker)

CRI-O Support in OpenShift

CRI-O tracks and versions identical to Kubernetes, simplifying support permutations



podman



A docker-compatible CLI
for containers

- Remote
management API
via Varlink
- Image/container
tagging
- Advanced
namespace
isolation

buildah



buildah

Secure & flexible OCI container builds

- Integrated into OCP build pods
- Performance improvements for knative enablement
- Image signing improvements

OpenShift 4 installation

Installer and
user-provisioned
infrastructure, bootstrap,
and more

OpenShift Bootstrap Process: Self-Managed Kubernetes

How to boot a self-managed cluster:

- OpenShift 4 is unique in that management extends all the way down to the operating system
- Every machine boots with a configuration that references resources hosted in the cluster it joins, enabling cluster to manage itself
- Downside is that every machine looking to join the cluster is waiting on the cluster to be created
- Dependency loop is broken using a bootstrap machine, which acts as a temporary control plane whose sole purpose is bringing up the permanent control plane nodes
- Permanent control plane nodes get booted and join the cluster leveraging the control plane on the bootstrap machine
- Once the pivot to the permanent control plane takes place, the remaining worker nodes can be booted and join the cluster

Bootstrapping process step by step:

1. Bootstrap machine boots and starts hosting the remote resources required for master machines to boot.
2. Master machines fetch the remote resources from the bootstrap machine and finish booting.
3. Master machines use the bootstrap node to form an etcd cluster.
4. Bootstrap node starts a temporary Kubernetes control plane using the newly-created etcd cluster.
5. Temporary control plane schedules the production control plane to the master machines.
6. Temporary control plane shuts down, yielding to the production control plane.
7. Bootstrap node injects OpenShift-specific components into the newly formed control plane.
8. Installer then tears down the bootstrap node or if user-provisioned, this needs to be performed by the administrator.

How everything deployed comes under management

Masters (Special)

- Terraform provisions initial masters*
- Machine API adopts existing masters post-provision
- Each master is a standalone Machine object
- Termination protection (avoid self-destruction)

Workers

- Each Machine Pool corresponds to MachineSet
- Optionally autoscale (min,max) and health check (replace if not ready > X minutes)

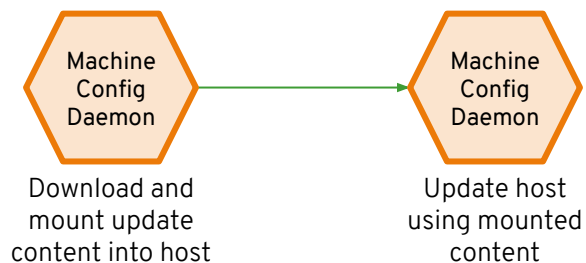
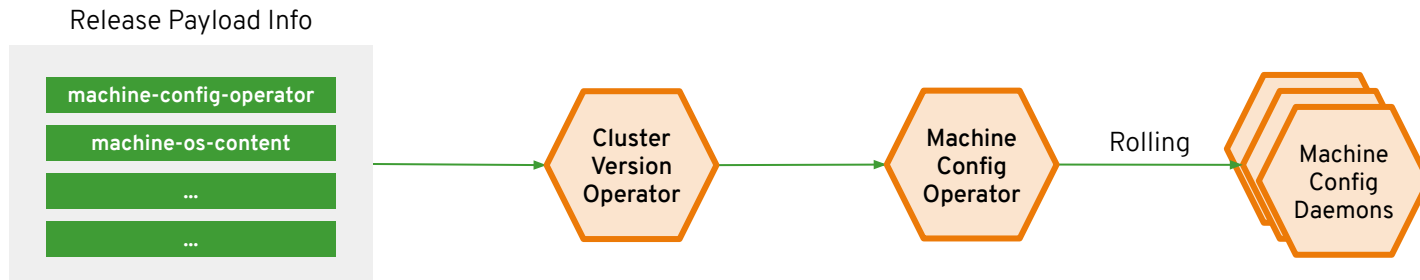
Multi-AZ

- MachineSets scoped to single AZ
- Installer stripes N machine sets across AZs by default
- Post-install best effort balance via cluster autoscaler

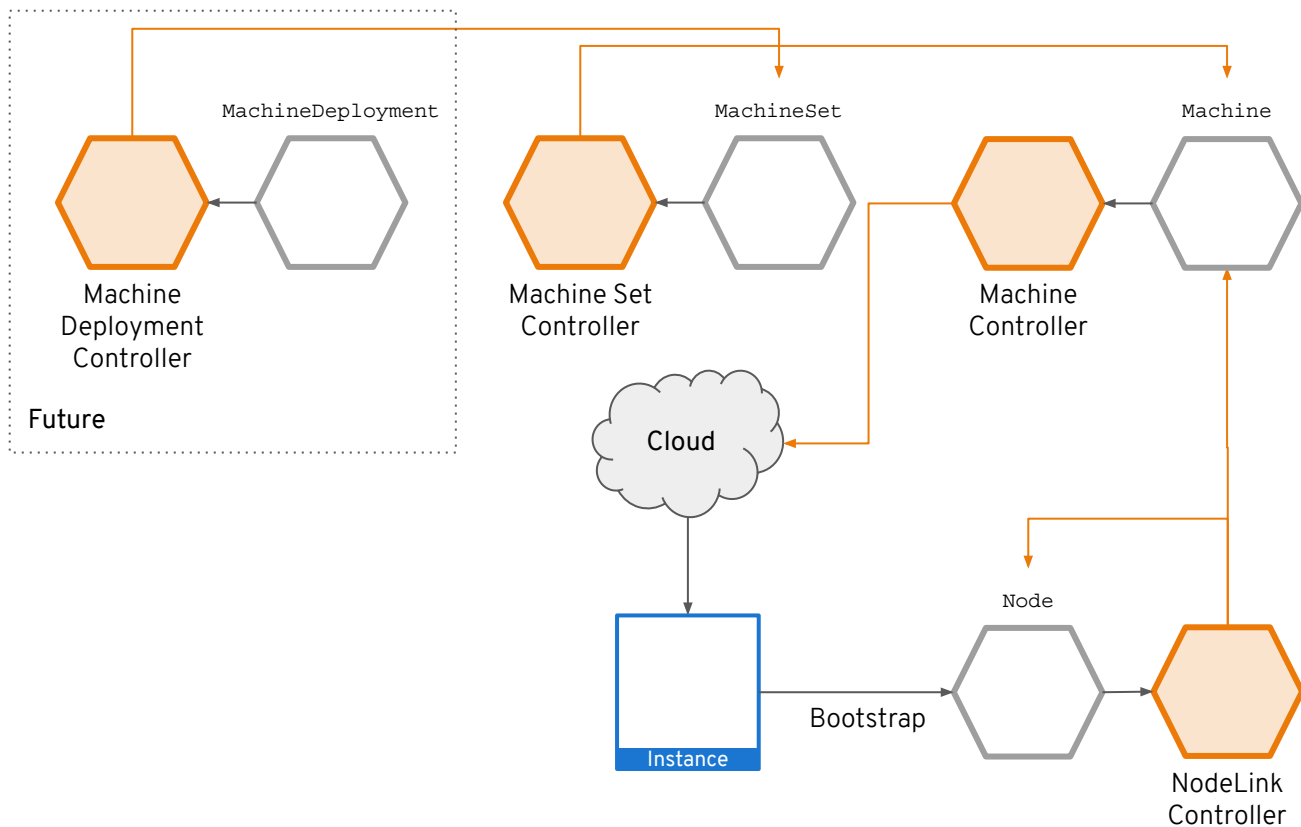
OpenShift 4 Cluster Management

Powered by Operators,
OpenShift 4 automates
many cluster
management activities

Over-the-air updates



Cloud API



OpenShift Security

Features, mechanisms
and processes for
container and platform
isolation



CONTROL

Application
Security

Container Content

CI/CD Pipeline

Container Registry

Deployment Policies



DEFEND

Infrastructure

Container Platform

Container Host Multi-tenancy

Network Isolation

Storage

Audit & Logging

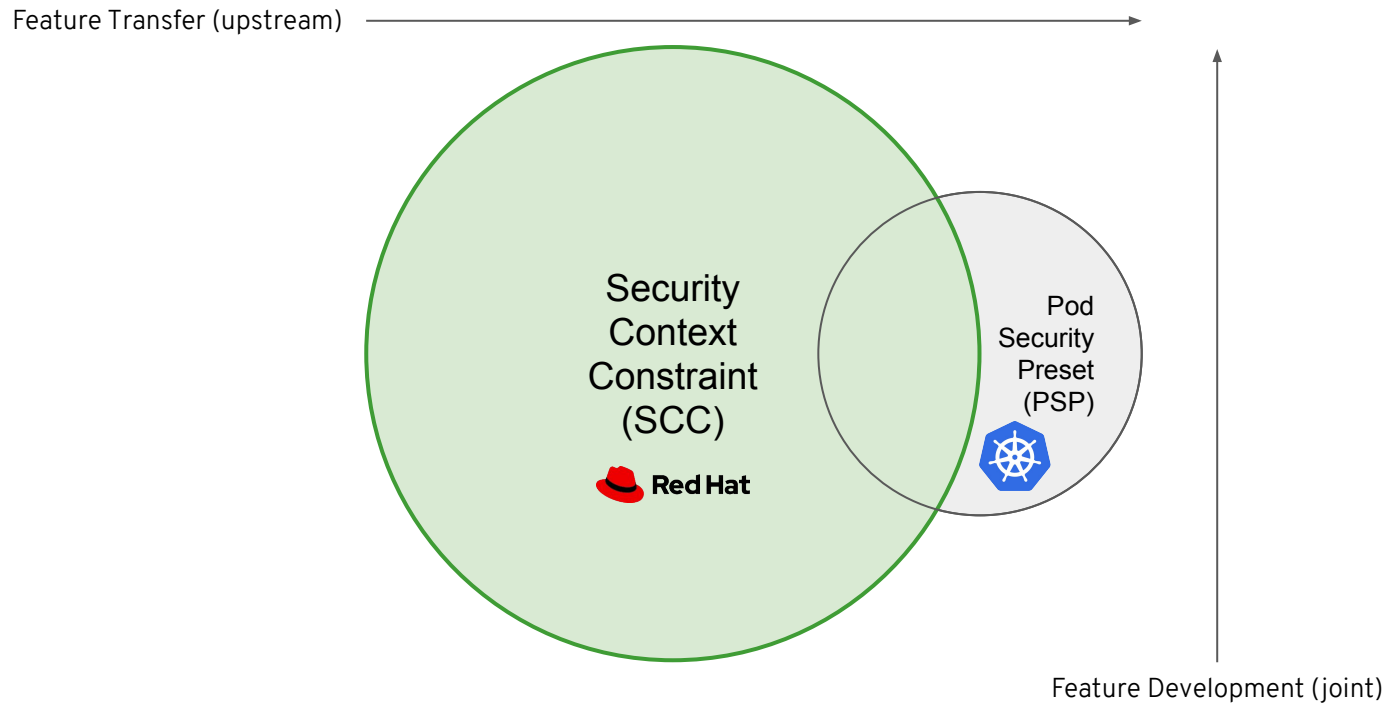
API Management



EXTEND

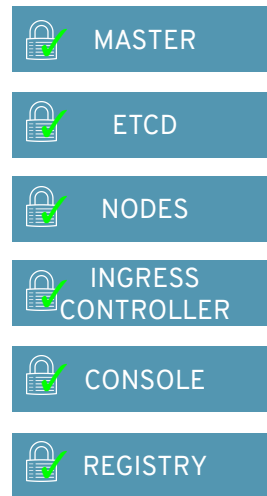
Security Ecosystem

Extended Depth of Protection

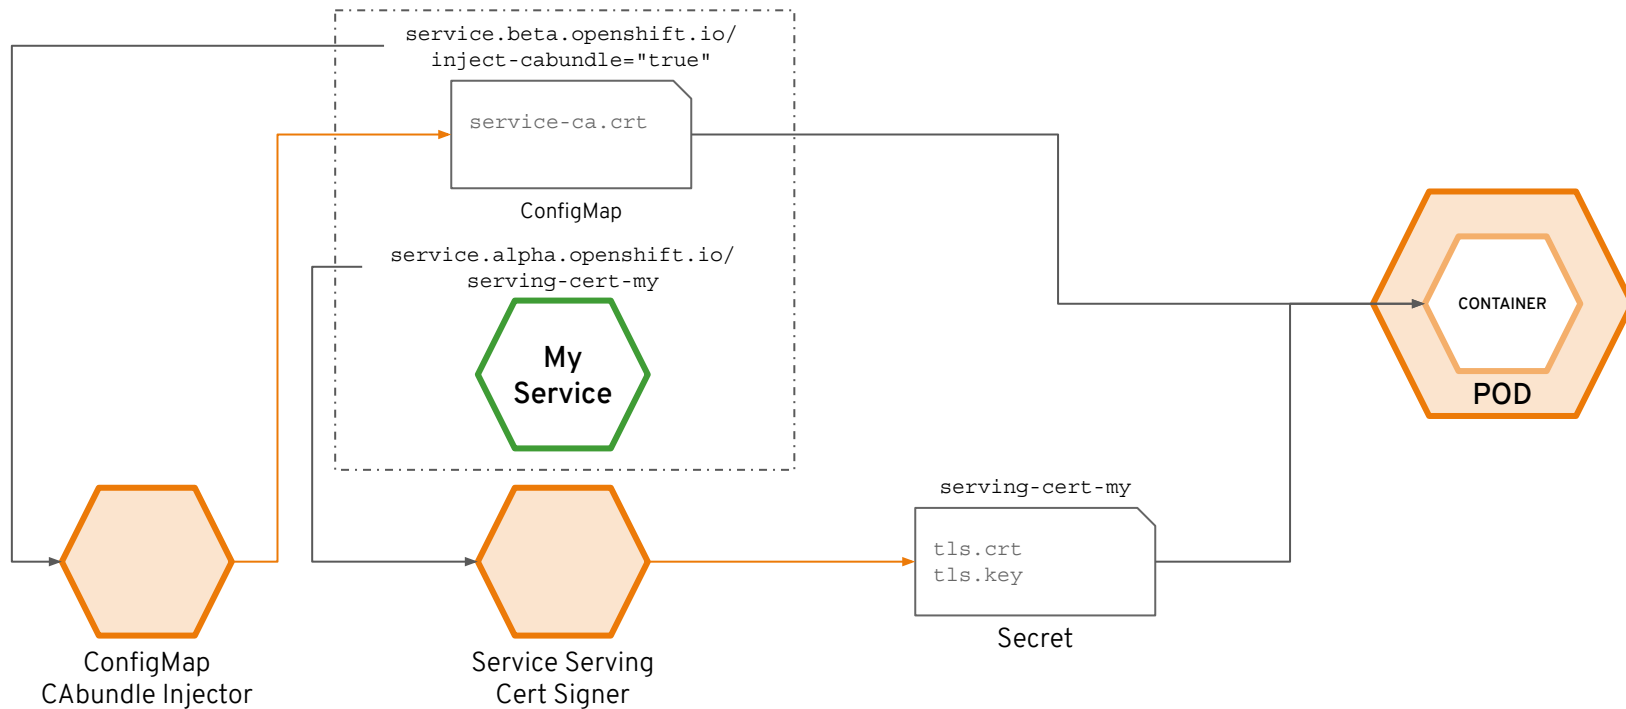


Certificates and Certificate Management

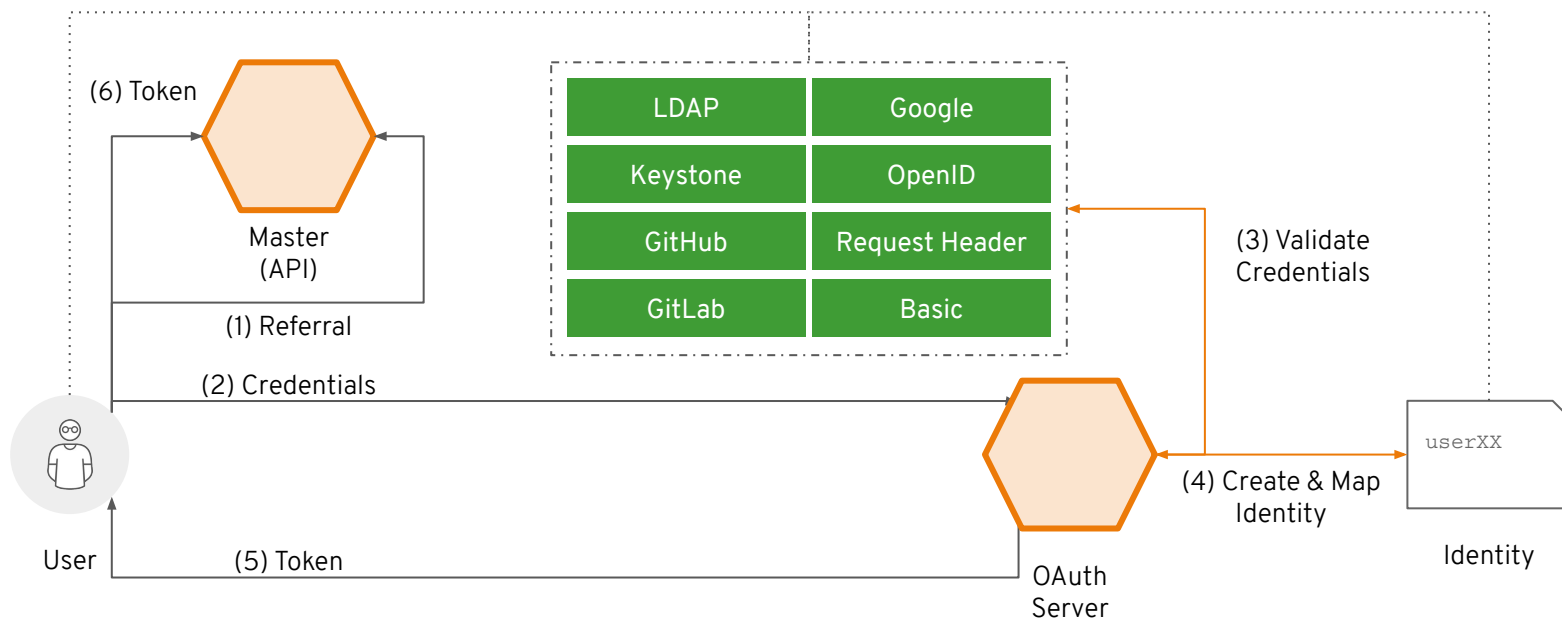
- OpenShift provides its own internal CA
- Certificates are used to provide secure connections to
 - master (APIs) and nodes
 - Ingress controller and registry
 - etcd
- Certificate rotation is automated
- Optionally configure external endpoints to use custom certificates



Service Certificates



Identity and Access Management



Fine-Grained RBAC

- Project scope & cluster scope available
- Matches request attributes (verb,object,etc)
- If no roles match, request is denied (deny by default)
- Operator- and user-level roles are defined by default
- Custom roles are supported

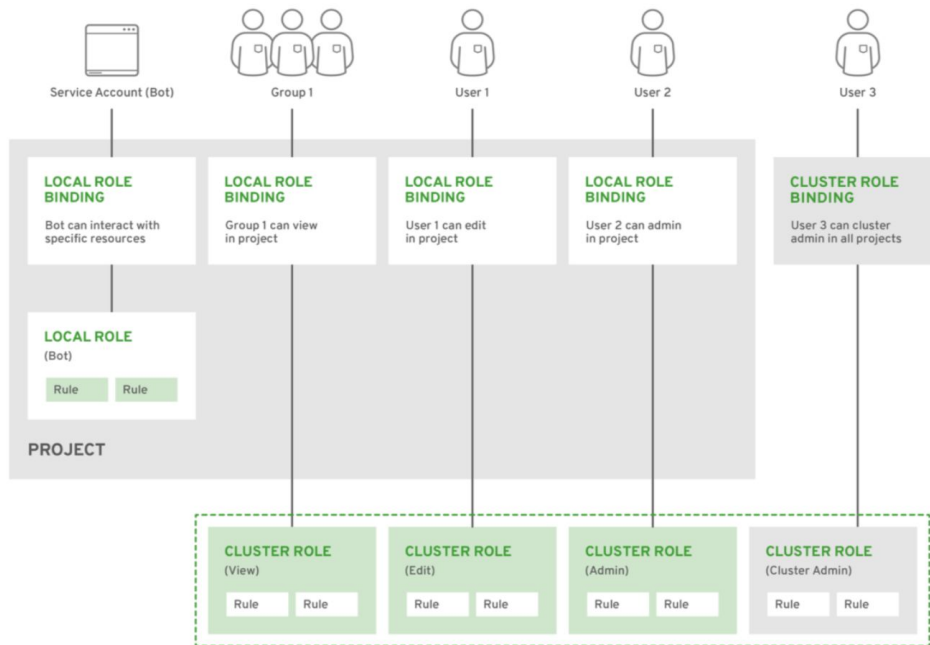


Figure 12 - Authorization Relationships

OpenShift Monitoring

An integrated cluster
monitoring and alerting
stack

OpenShift Cluster Monitoring



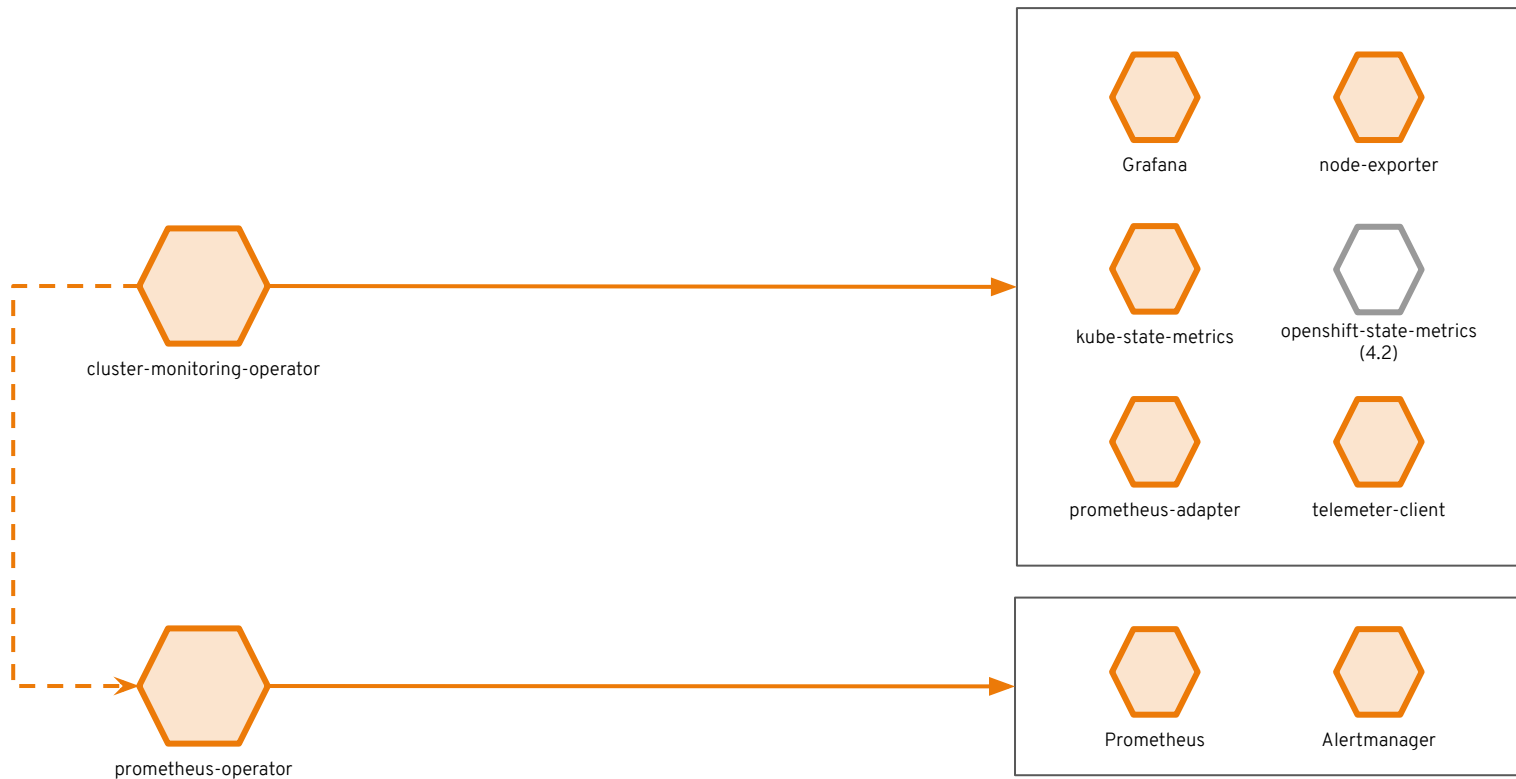
Metrics collection and storage
via Prometheus, an
open-source monitoring system
time series database.

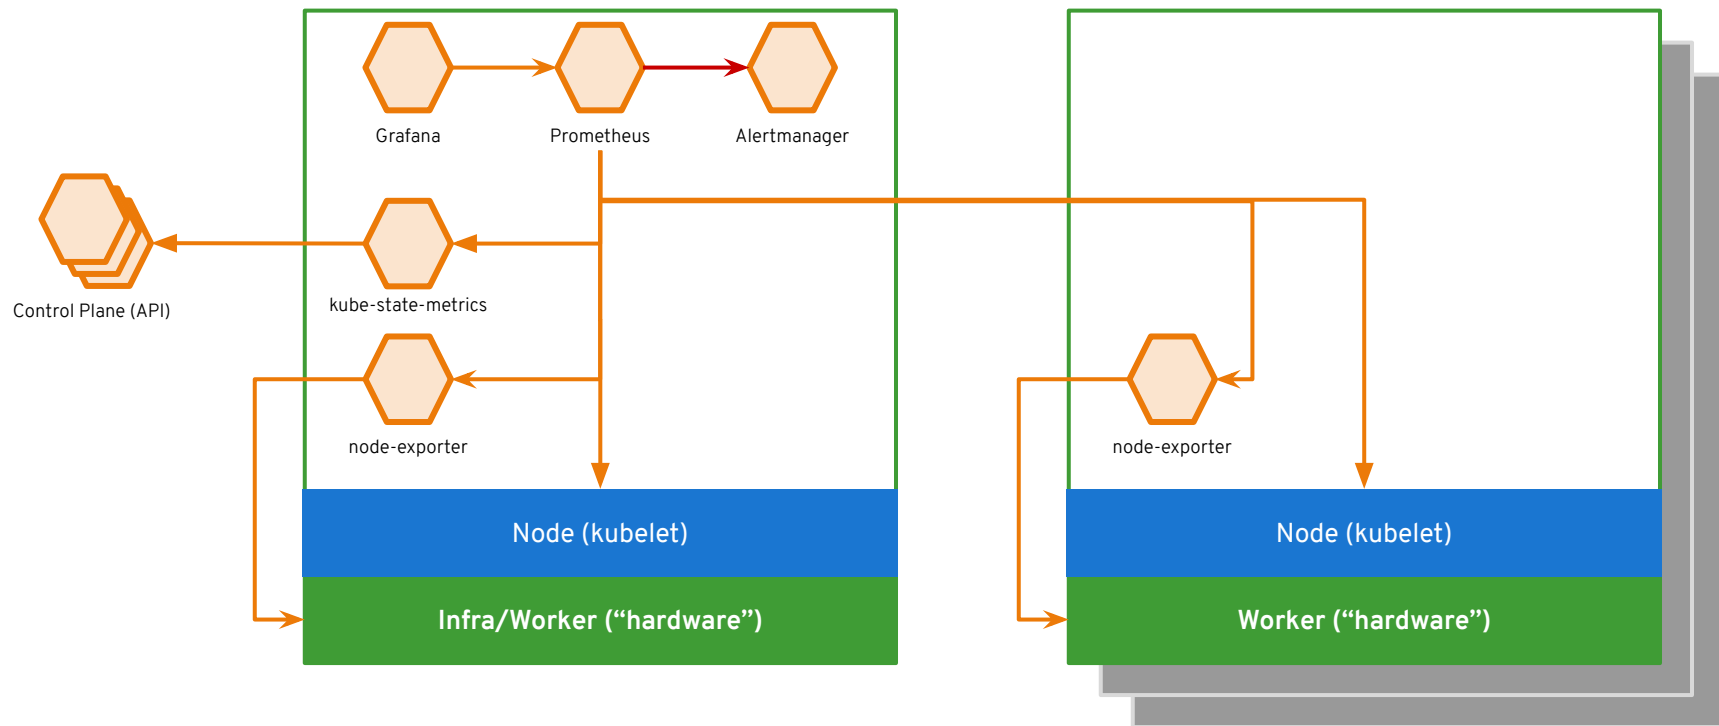


Alerting/notification via
Prometheus' Alertmanager, an
open-source tool that handles
alerts send by Prometheus.



Metrics visualization via
Grafana, the leading metrics
visualization technology.





OpenShift Logging

An integrated solution for
exploring and
corroborating application
logs

Observability via log exploration and corroboration with EFK

Components

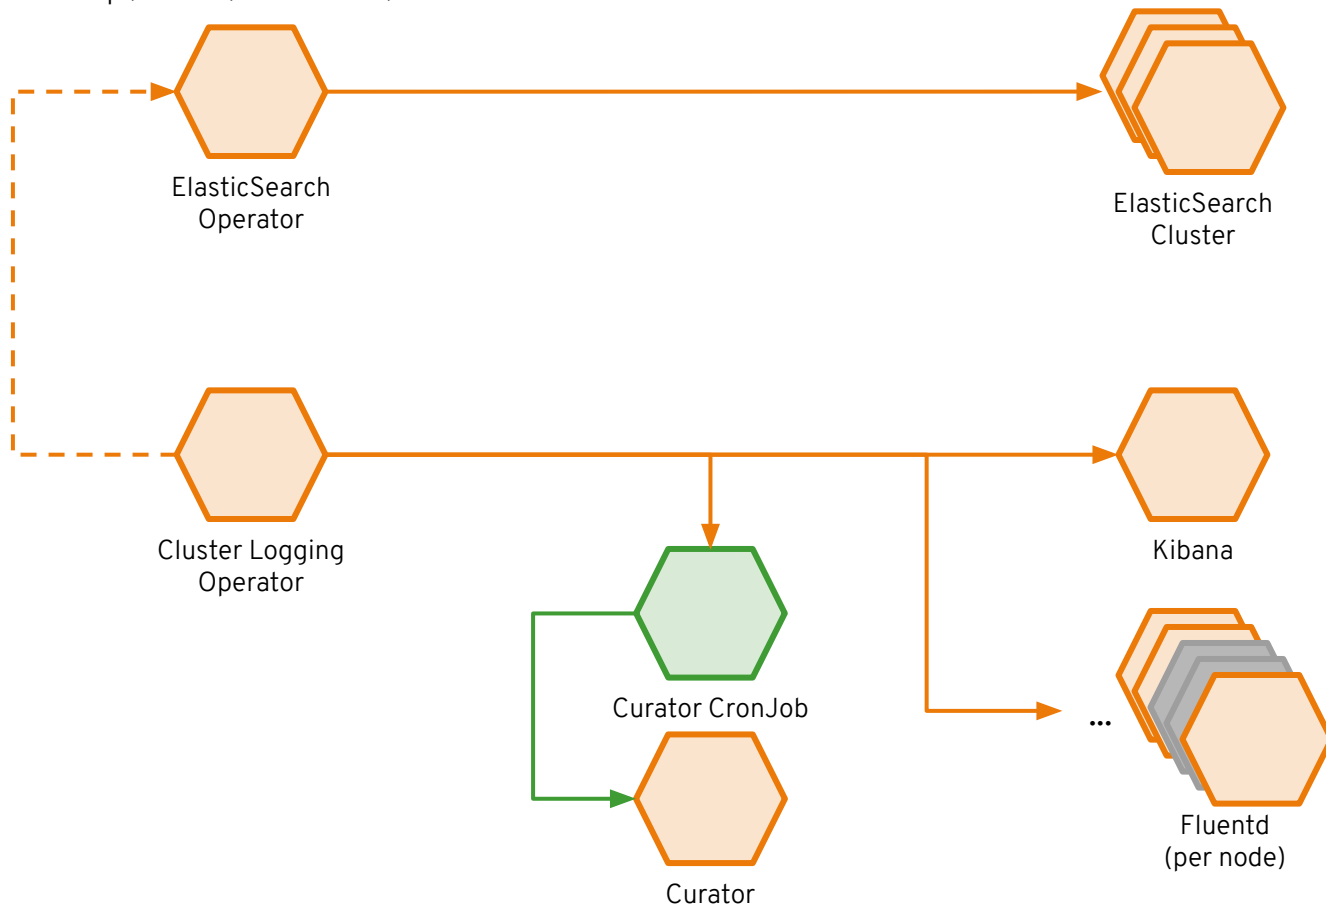
- **Elasticsearch:** a search and analytics engine to store logs
- **Fluentd:** gathers logs and sends to Elasticsearch.
- **Kibana:** A web UI for Elasticsearch.

Access control

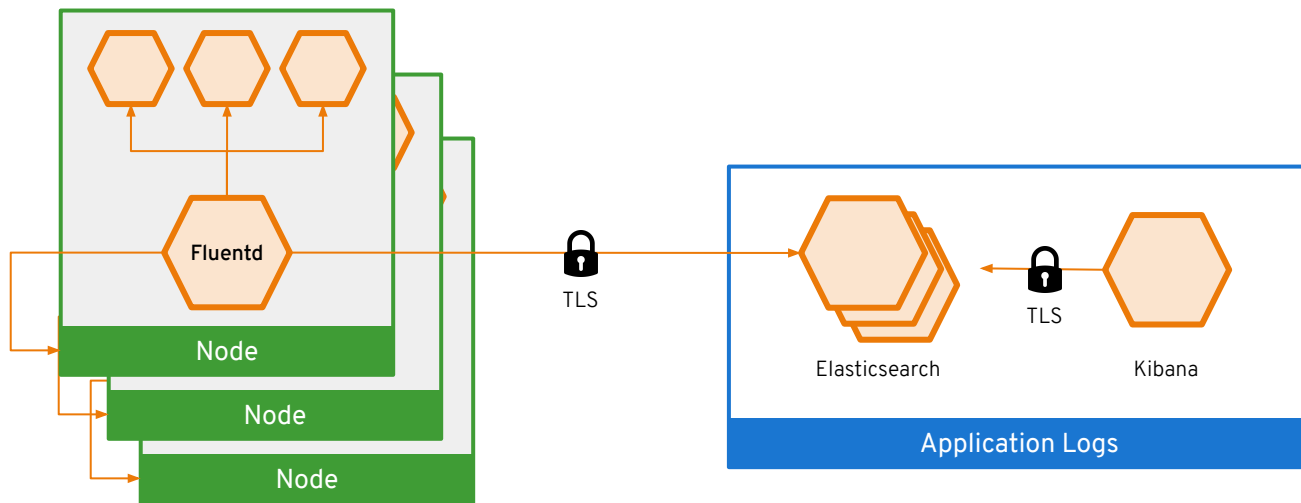
- Cluster administrators can view all logs
- Users can only view logs for their projects

Ability to forward logs elsewhere

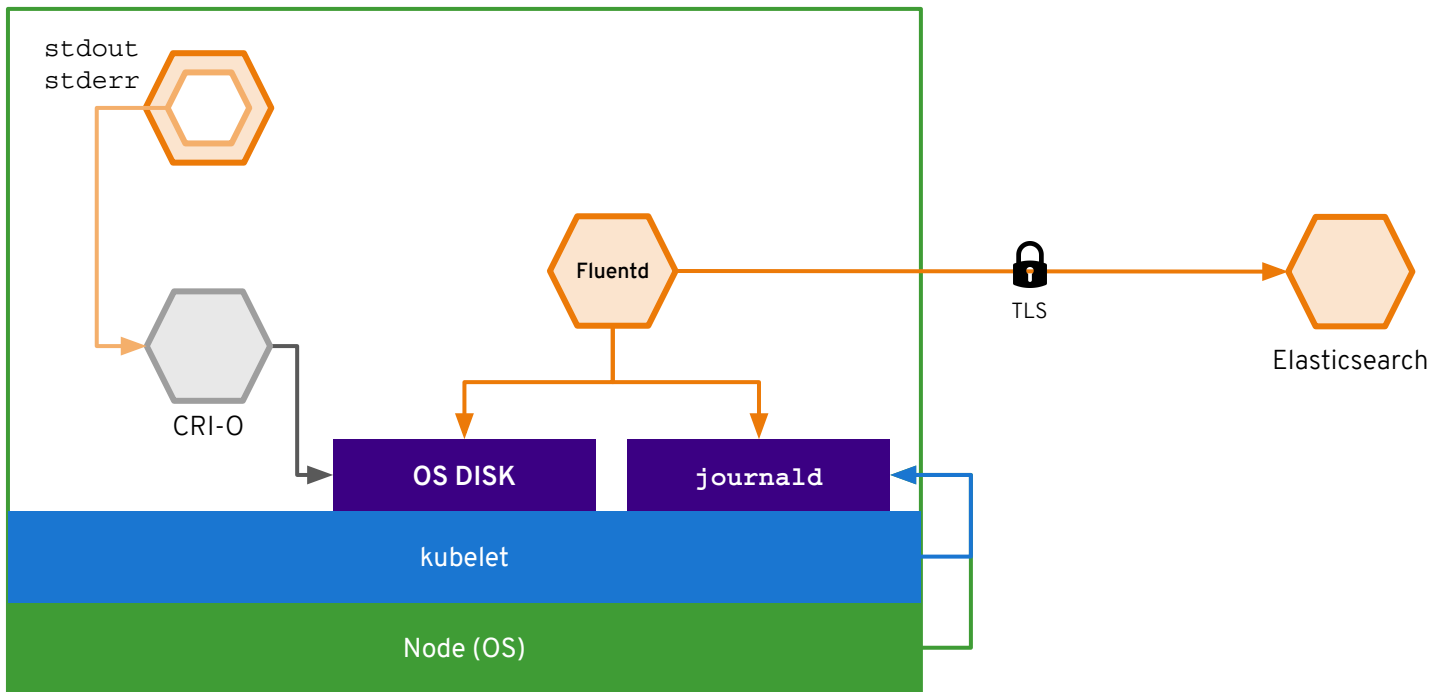
- External elasticsearch, Splunk, etc



Log data flow in OpenShift



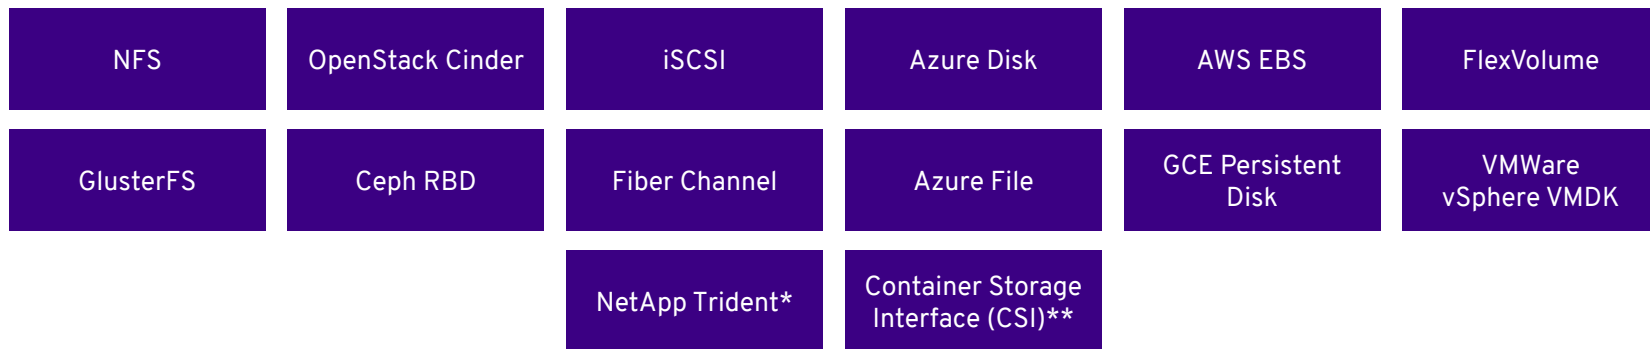
Log data flow in OpenShift



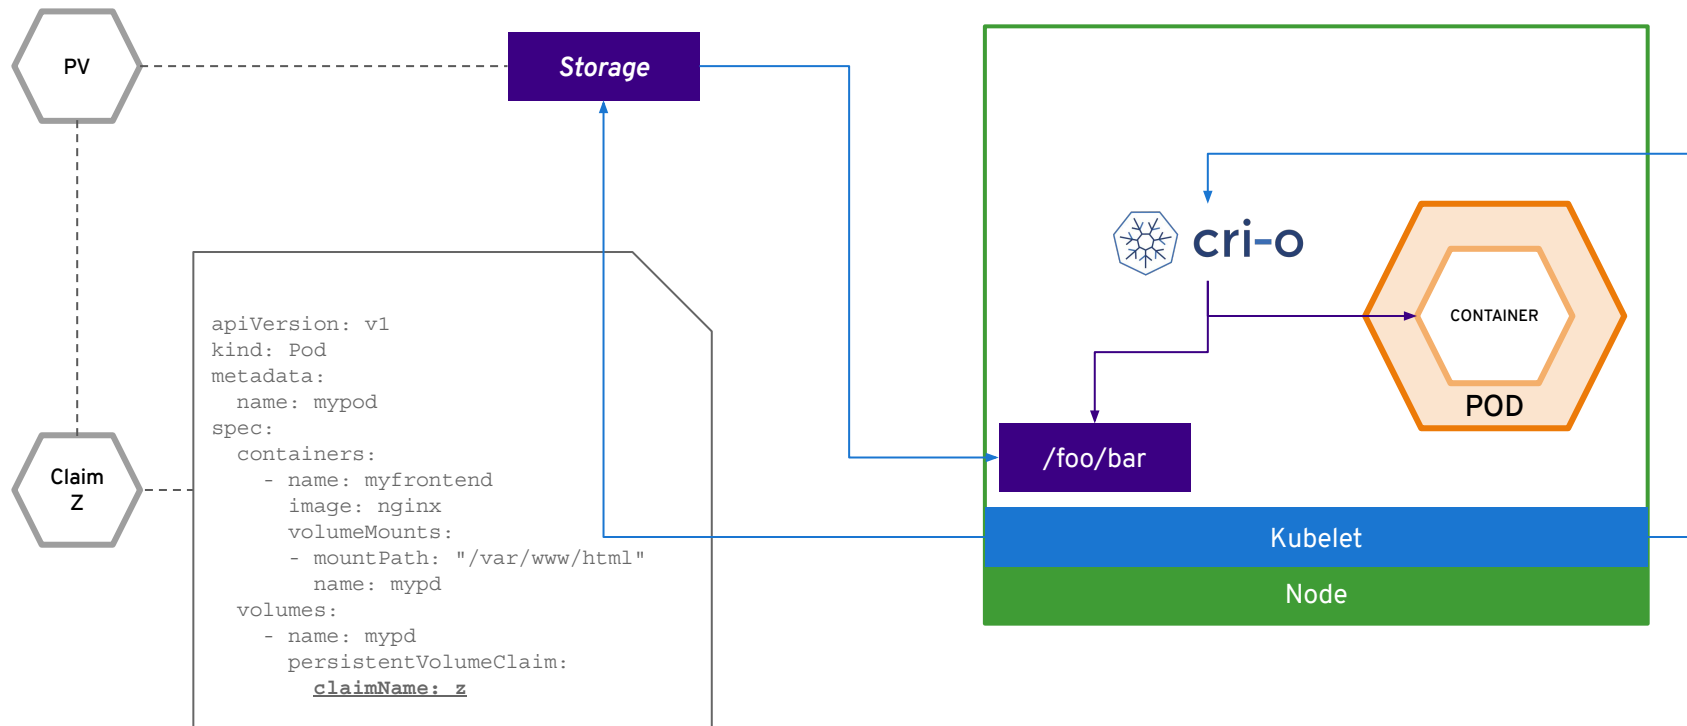
Persistent Storage

Connecting real-world
storage to your
containers to enable
stateful applications

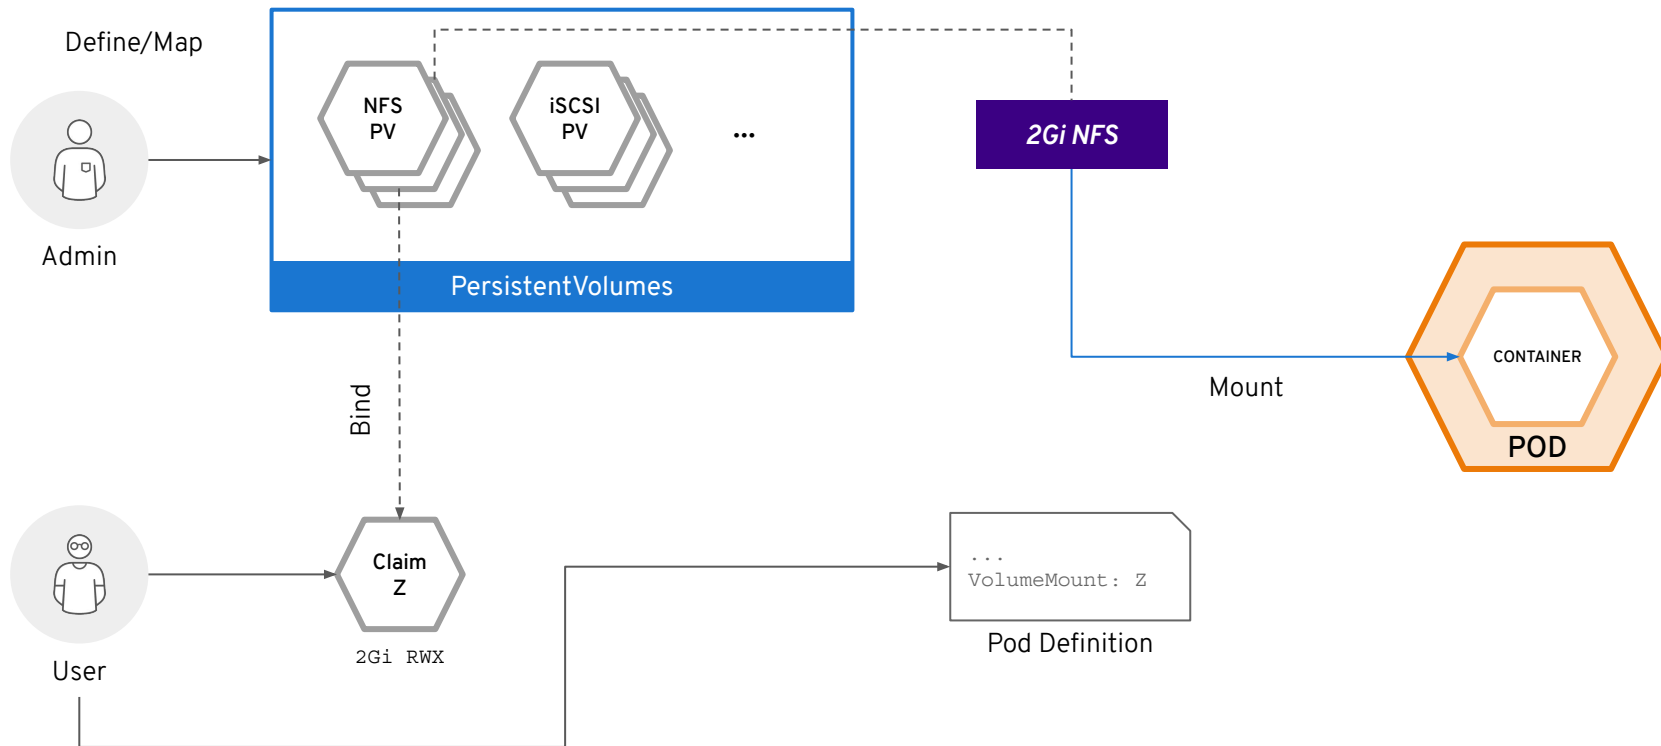
A broad spectrum of static and dynamic storage endpoints



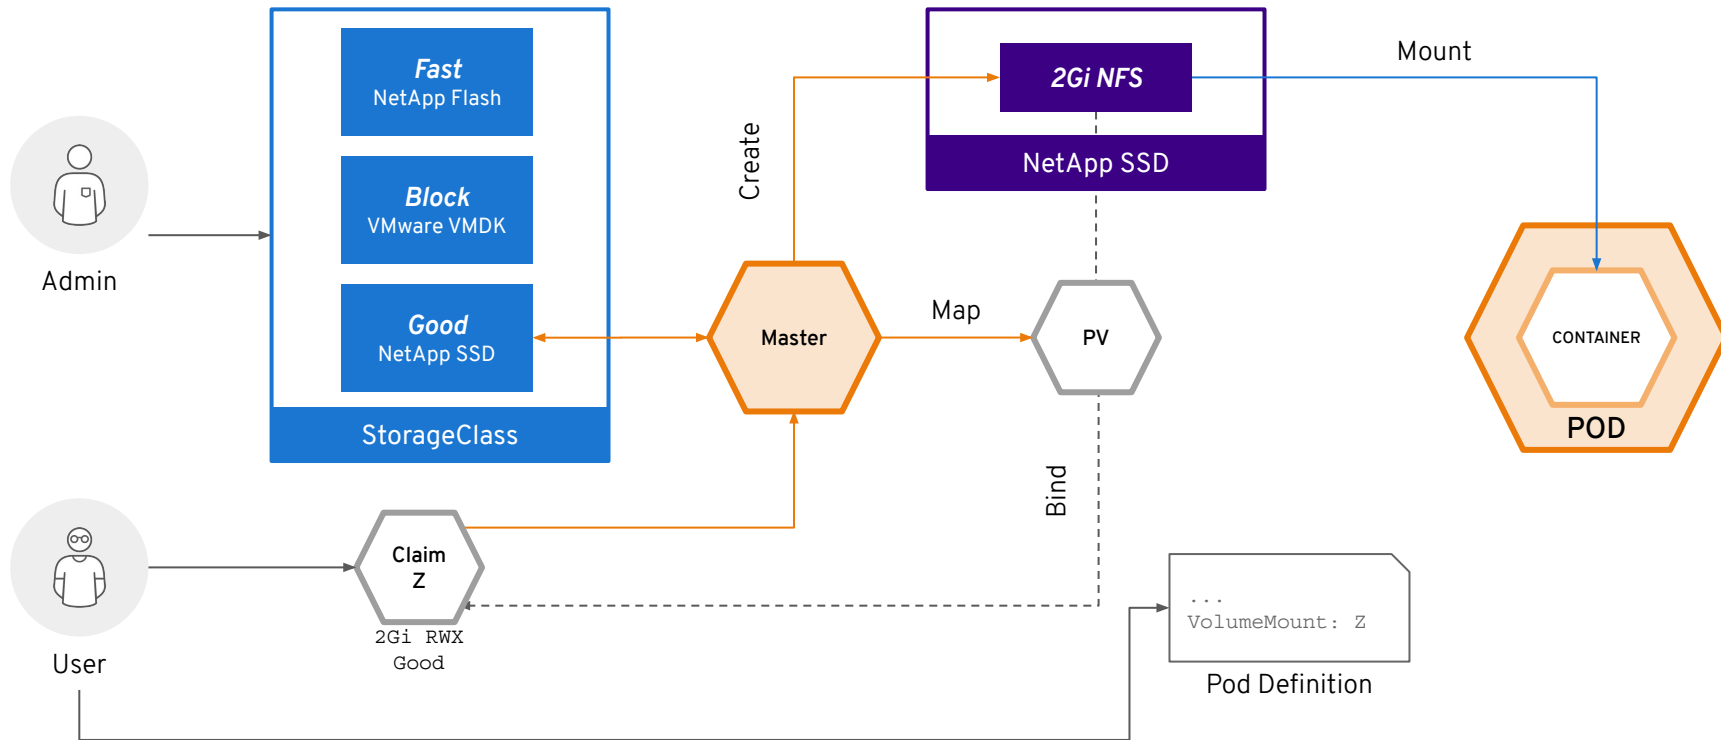
PV Consumption



Static Storage Provisioning

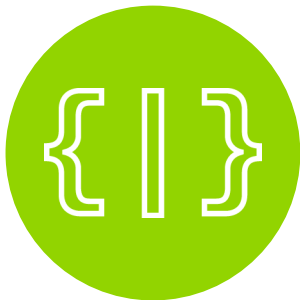


Dynamic Storage Provisioning



Build and Deploy Container Images

Tools and automation
that makes developers
productive quickly



**DEPLOY YOUR
SOURCE CODE**



**DEPLOY YOUR
APP BINARY**



**DEPLOY YOUR
CONTAINER IMAGE**

