



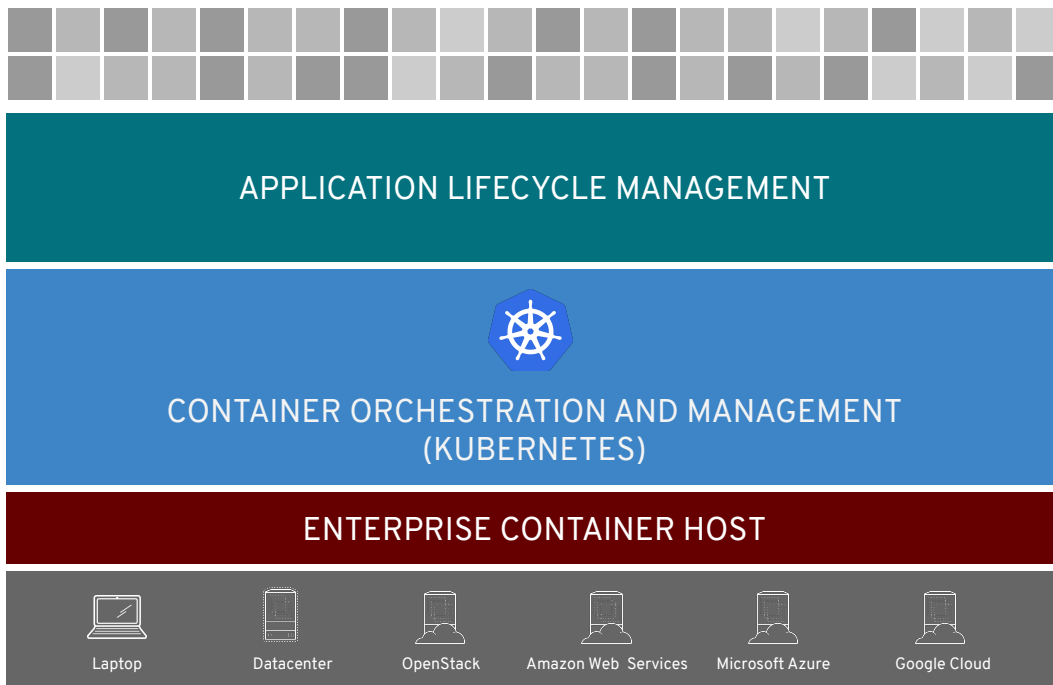
OpenShift 4.x Architecture Workshop

OpenShift Container Platform (OCP)
Advanced Architecture

July 2019




OPENSIFT CONTAINER PLATFORM

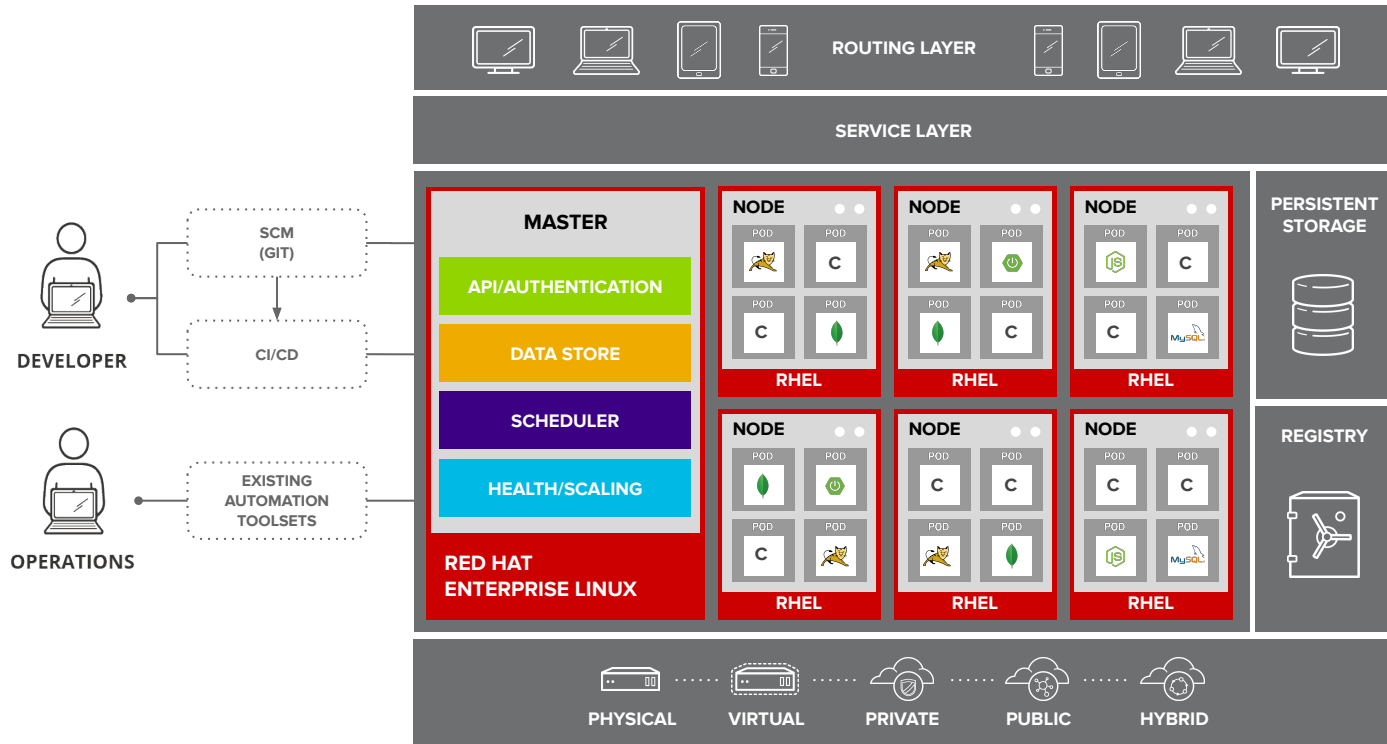


ANY
INFRASTRUCTURE

OPENSIFT CONTAINER PLATFORM



OPENSIFT ARCHITECTURE



Container Concepts Overview

A container is the smallest compute unit

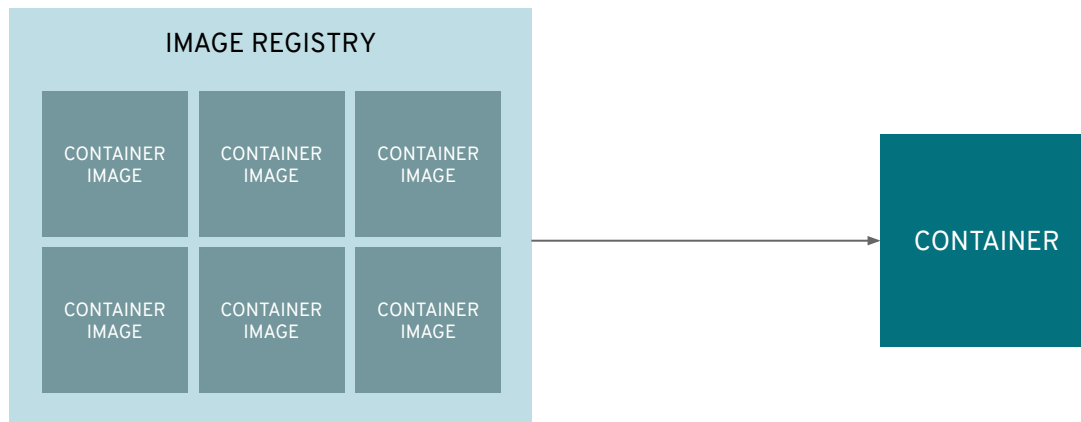


CONTAINER

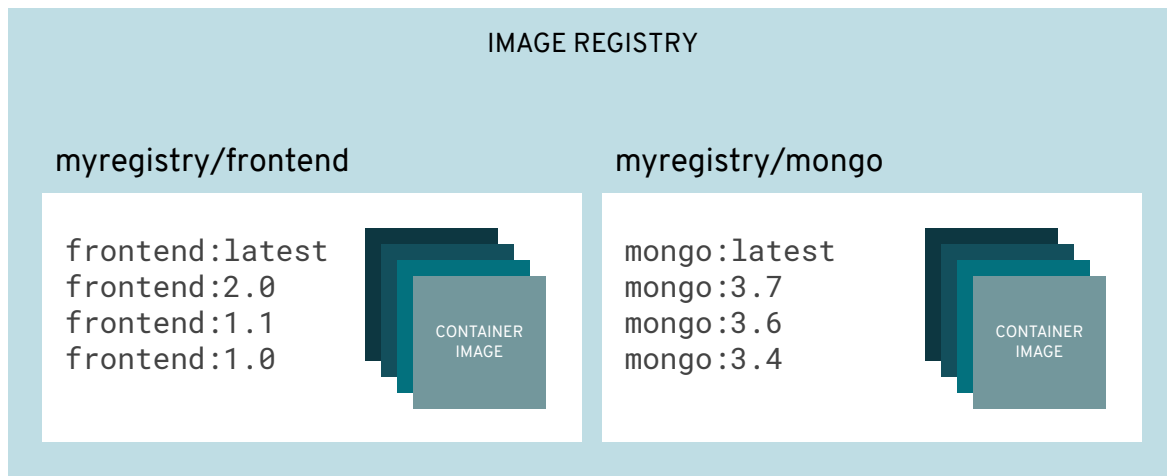
containers are created from
container images




container images are stored in an image registry




an image repository contains all versions of an image in the image registry



containers are wrapped in pods which are
units of deployment and management



Pods configuration is defined in a deployment

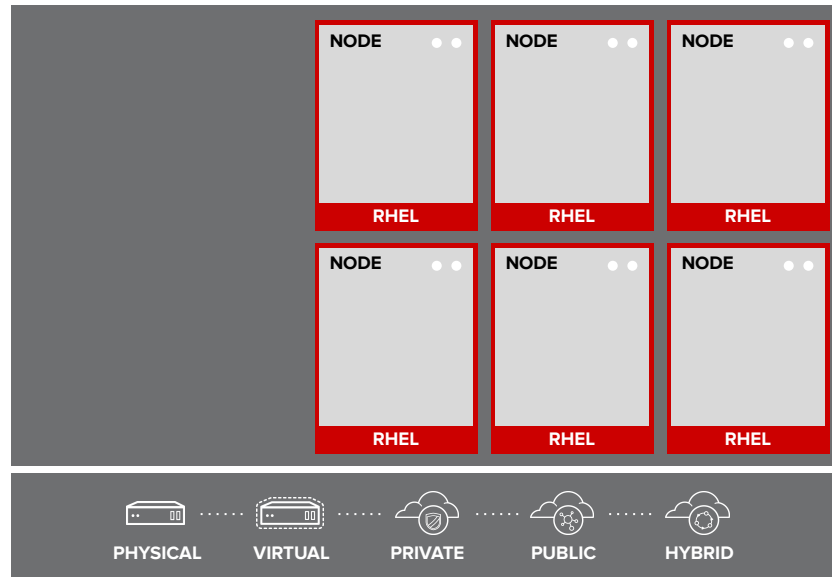


OpenShift Architecture


YOUR CHOICE OF INFRASTRUCTURE



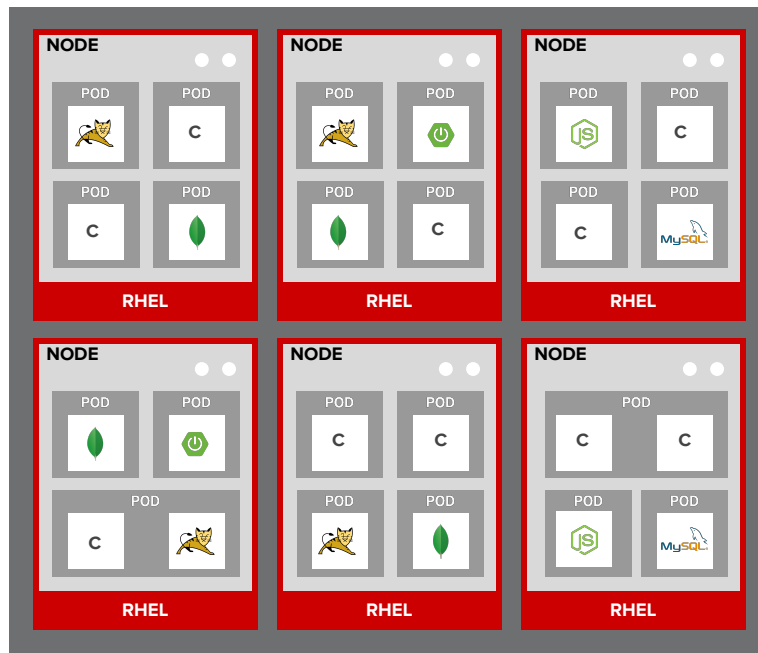
NODES RHEL INSTANCES WHERE APPS RUN



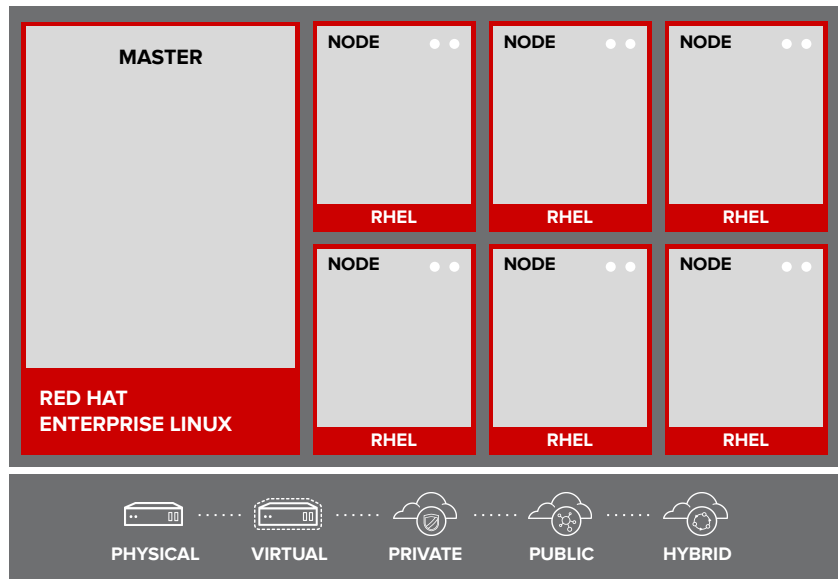
APPS RUN IN CONTAINERS



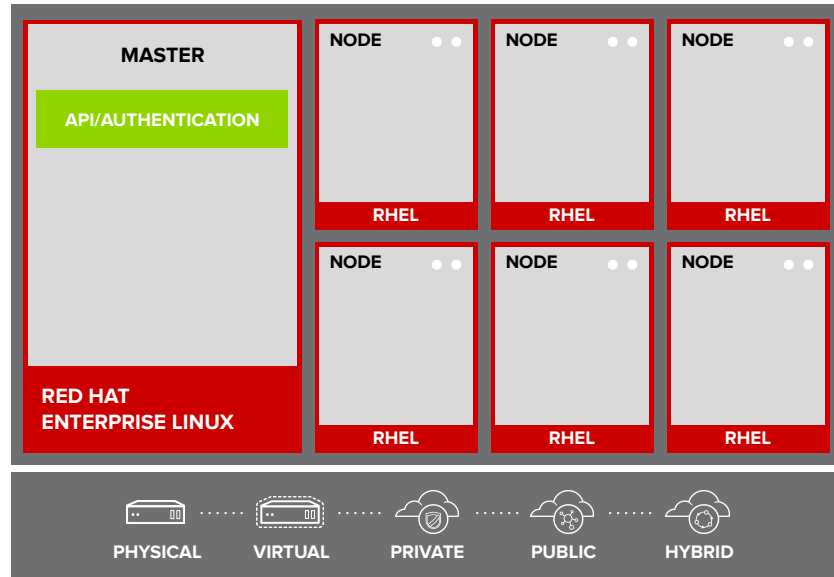
PODS ARE THE UNIT OF ORCHESTRATION



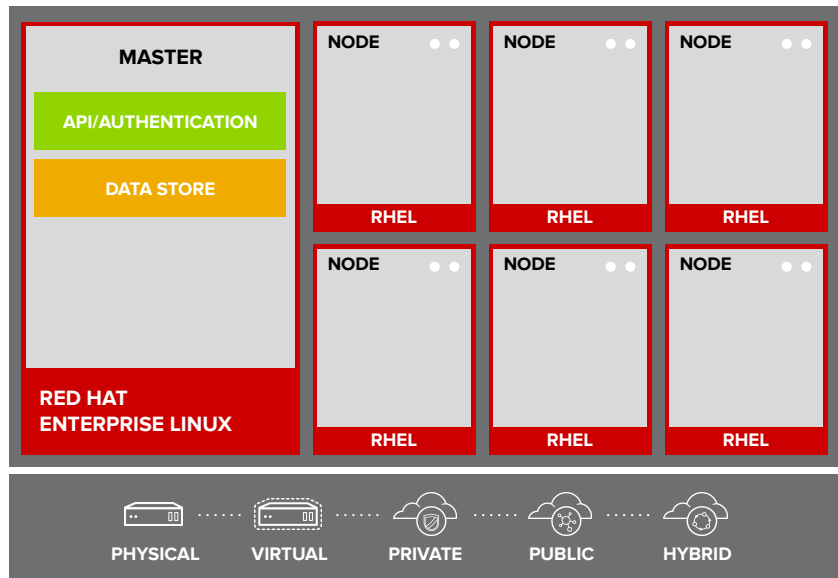
MASTERS ARE THE CONTROL PLANE



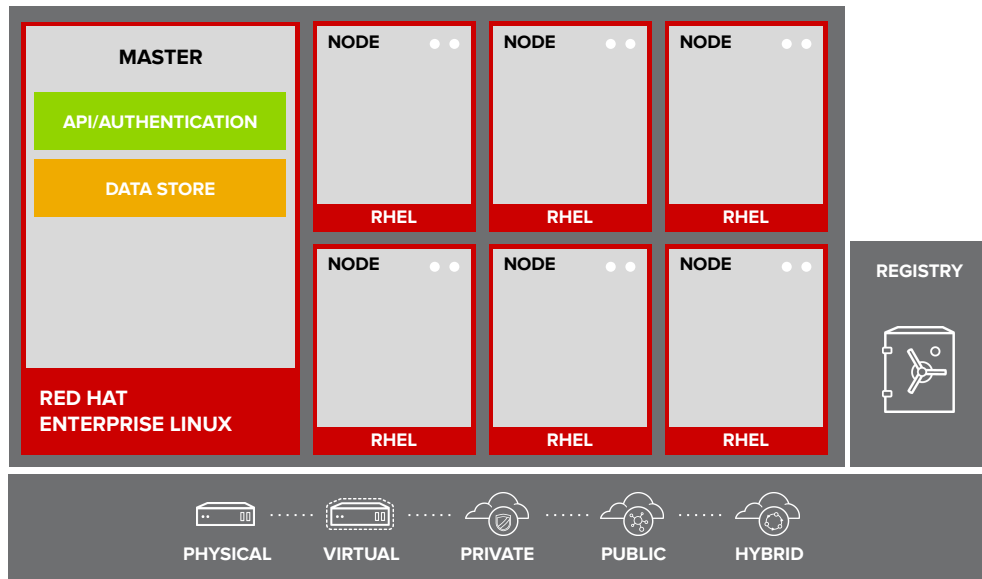
API AND AUTHENTICATION



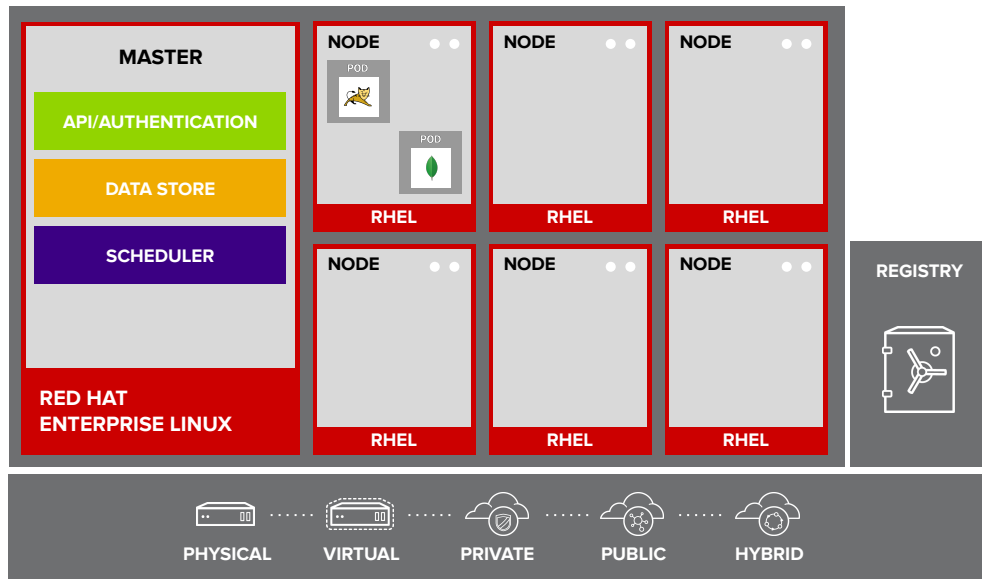
DESIRED AND CURRENT STATE



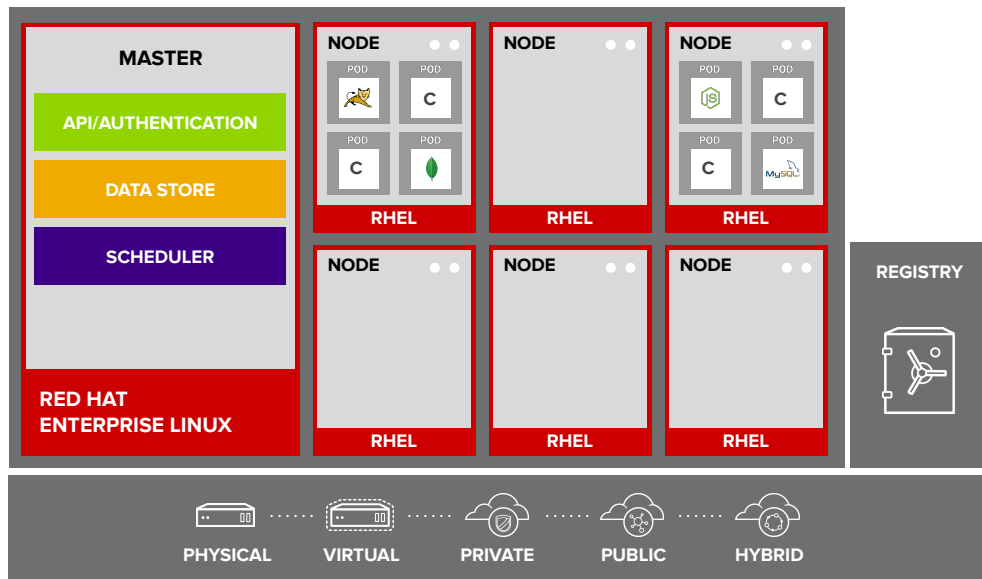
INTEGRATED CONTAINER REGISTRY



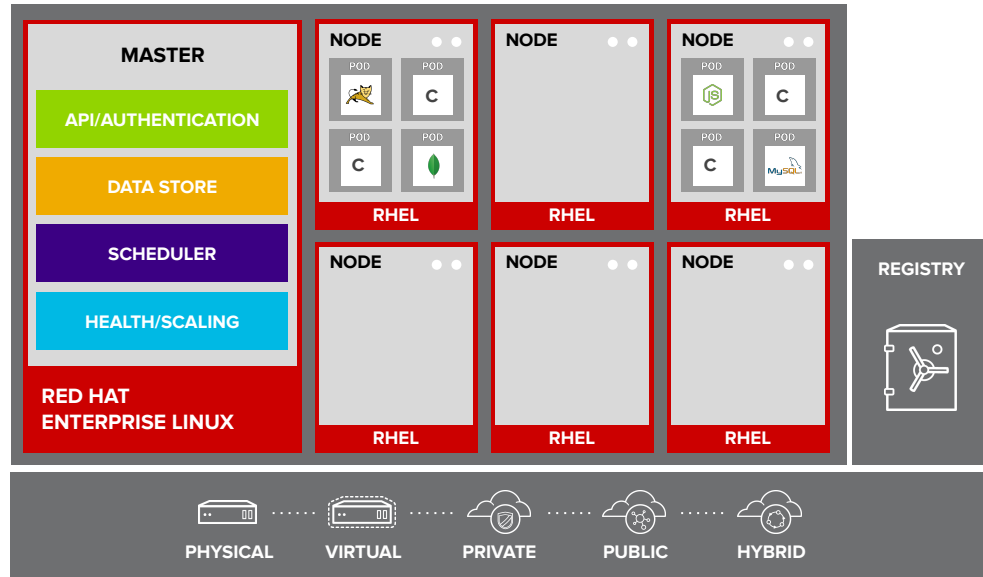
ORCHESTRATION AND SCHEDULING



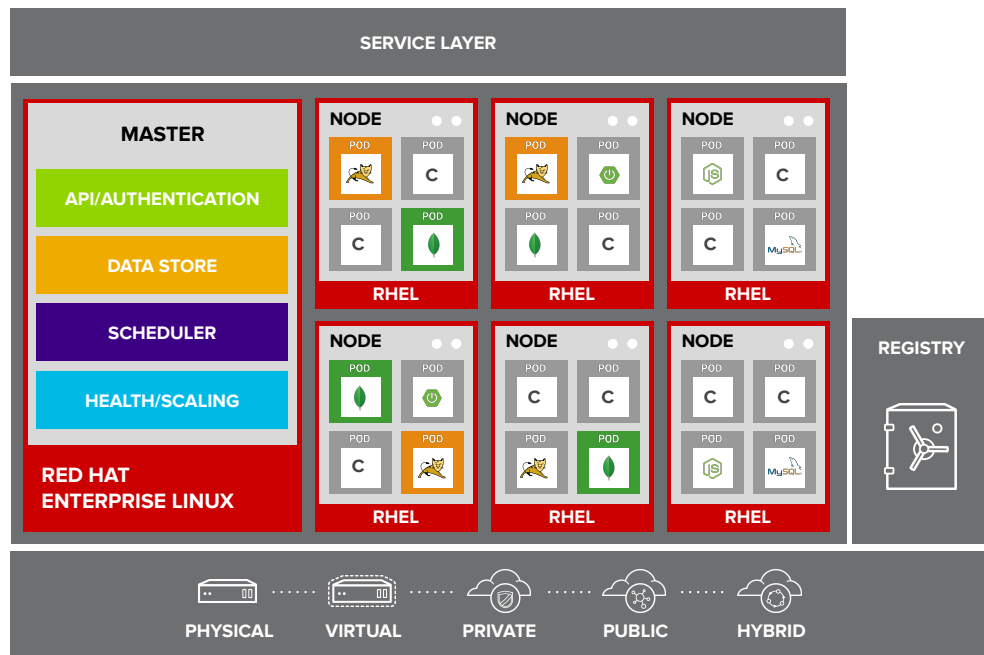
PLACEMENT BY POLICY



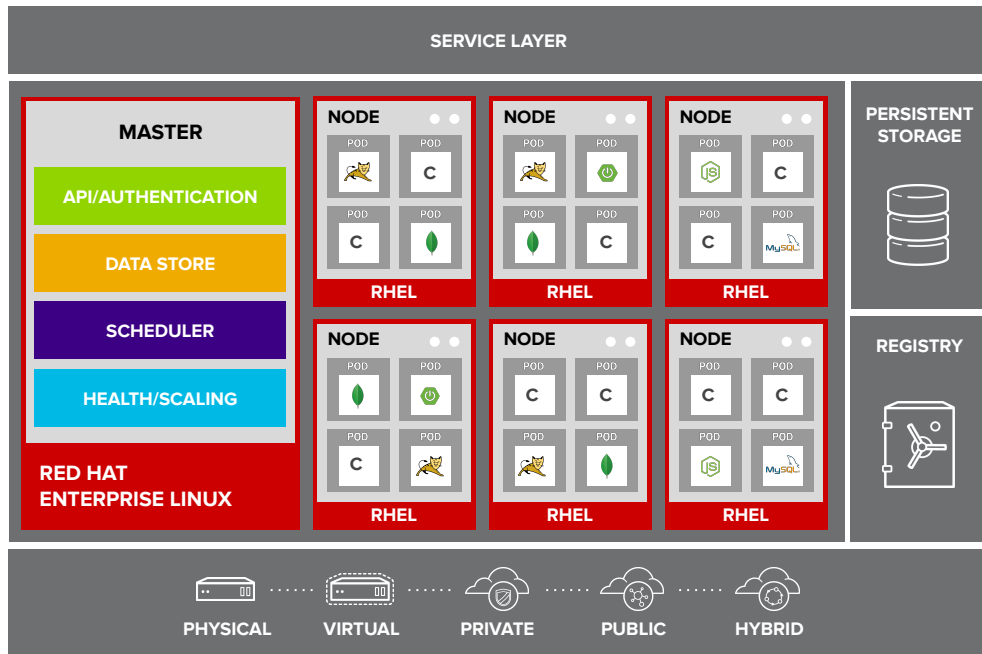
AUTOSCALING PODS



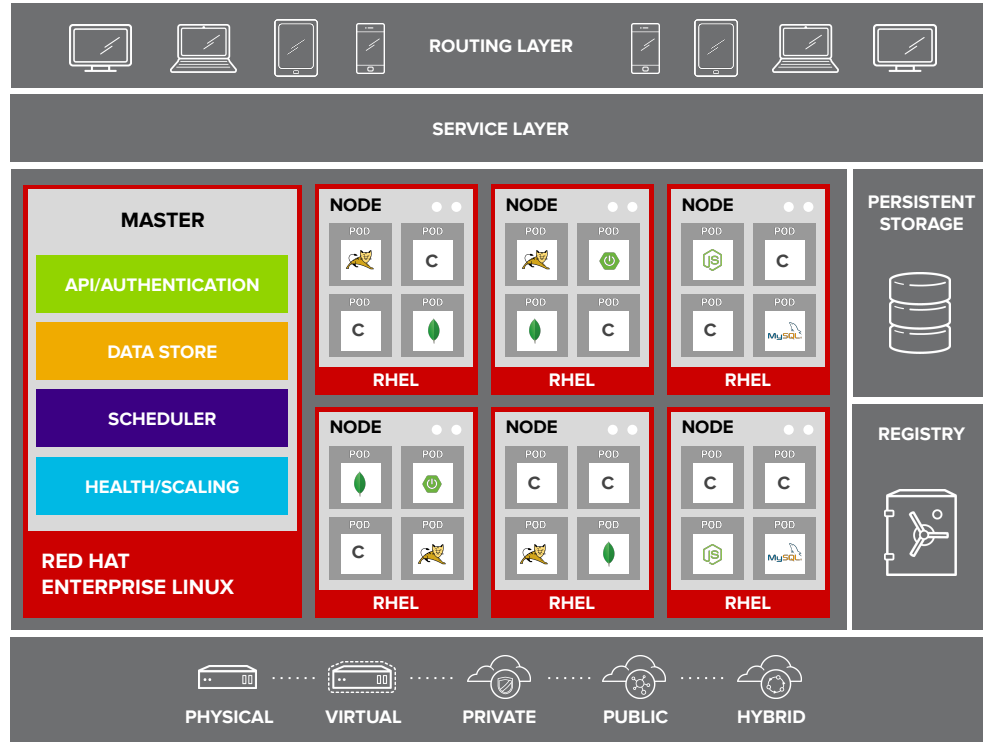
SERVICE DISCOVERY



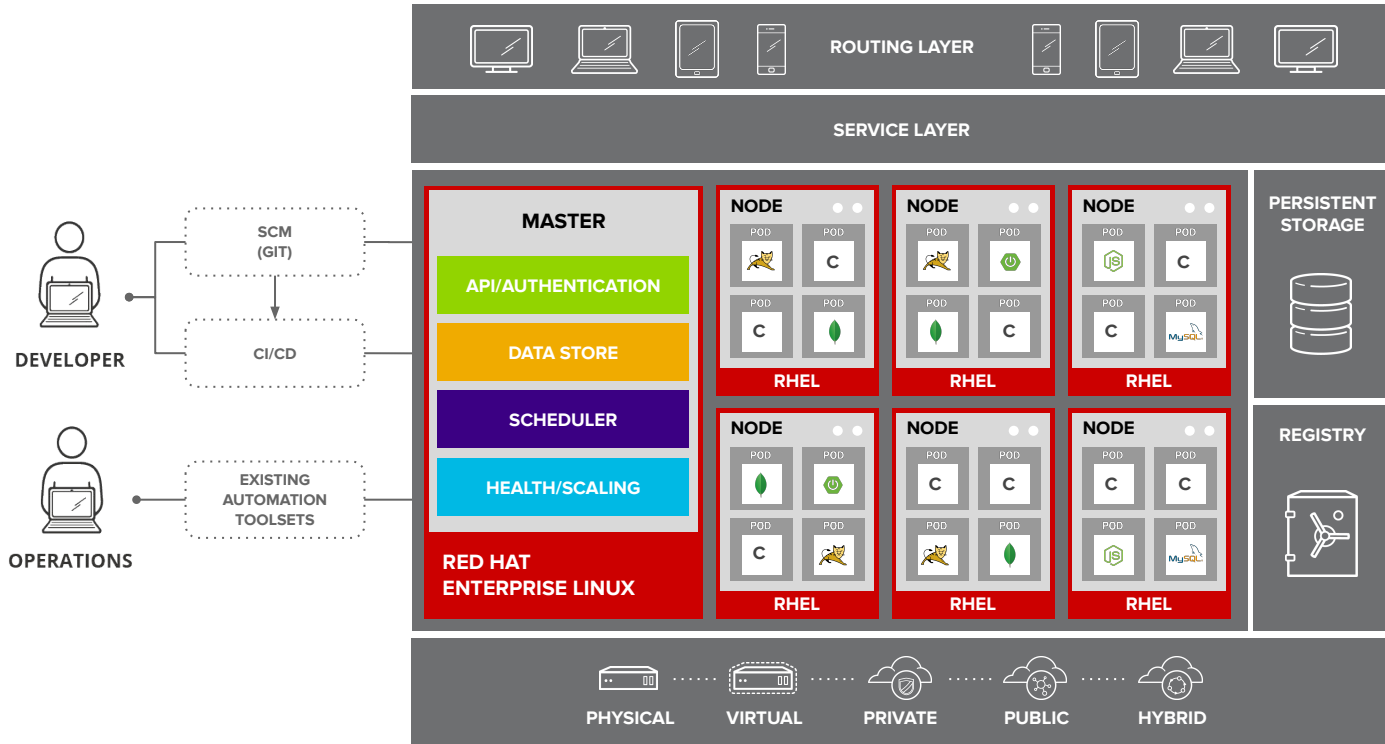
PERSISTENT DATA IN CONTAINERS



ROUTING AND LOAD-BALANCING




ACCESS VIA WEB, CLI, IDE AND API



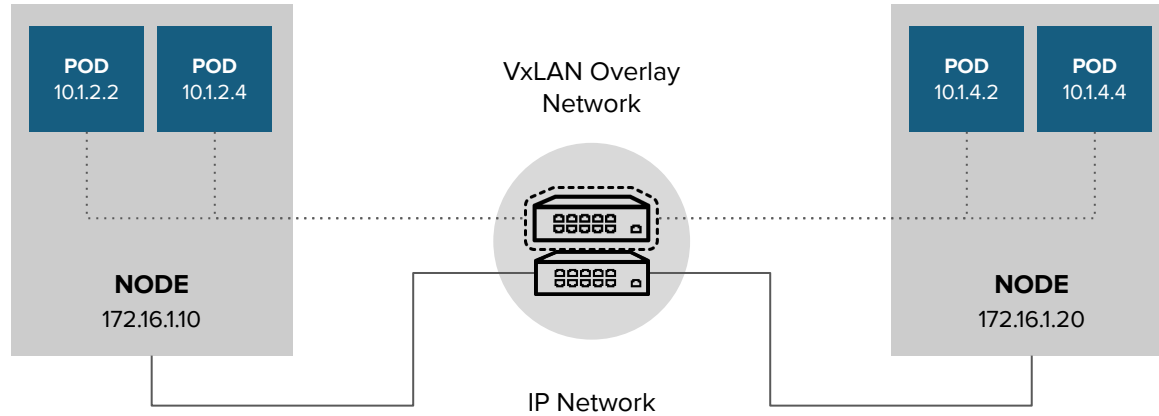
Networking

OPENSHIFT NETWORKING

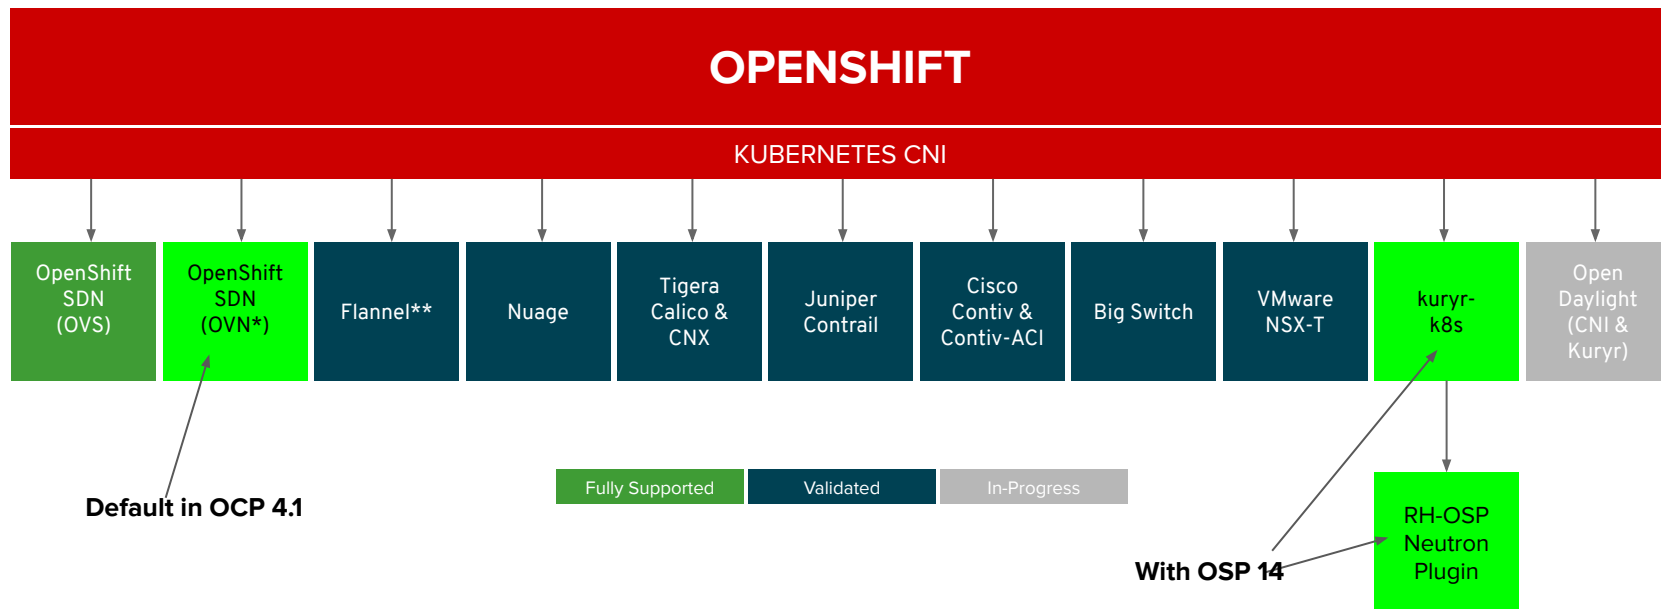
- Built-in internal DNS to reach services by name
- Split DNS is supported via DNSmasq
 - Master answers DNS queries for internal services
 - Other name servers serve the rest of the queries
- Software Defined Networking (SDN) for a unified cluster network to enable pod-to-pod communication
- OpenShift follows the Kubernetes Container Networking Interface (CNI) plug-in model



OPENSIFT NETWORKING



OPENSHIFT NETWORK PLUGINS



OPENSIFT SDN

FLAT NETWORK


- All pods can communicate with each other across projects

MULTI-TENANT NETWORK

- Project-level network isolation
- Multicast support
- Egress network policies

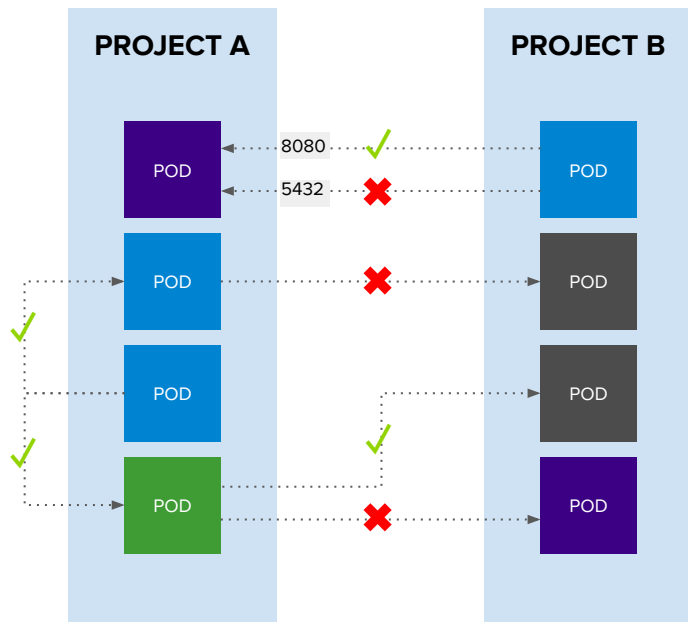
NETWORK POLICY (Default)

- Granular policy-based isolation



Multi-Tenant Network

OPENSIFT SDN - NETWORK POLICY

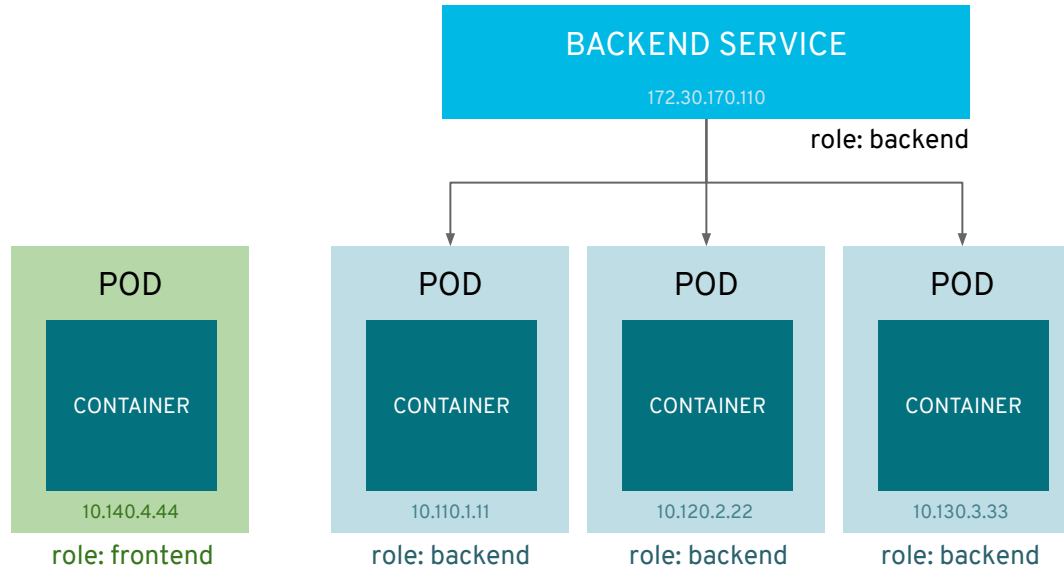


Example Policies

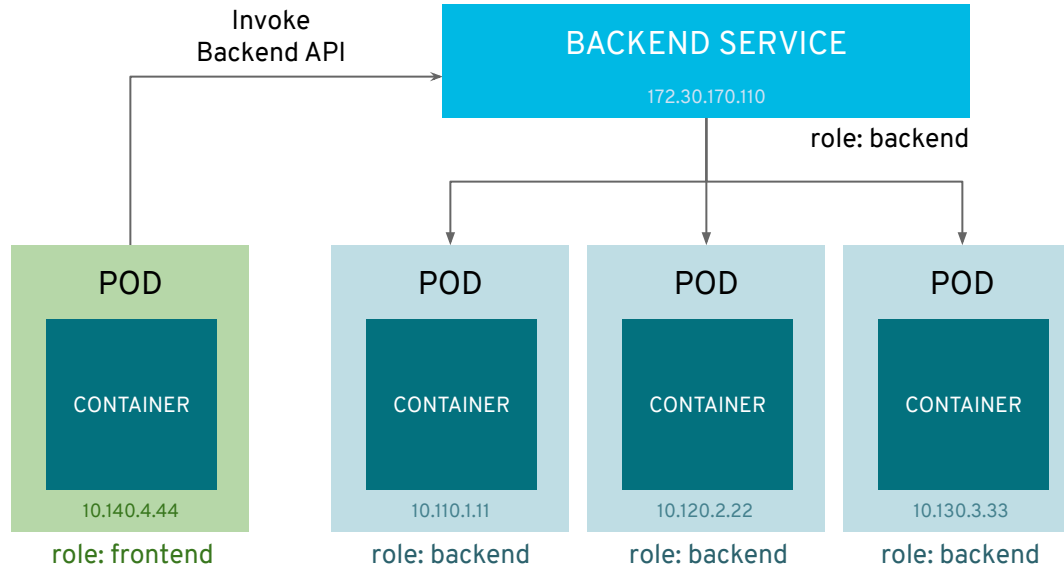
- Allow all traffic inside the project
- Allow traffic from green to gray
- Allow traffic to purple on 8080

```
apiVersion: extensions/v1beta1
kind: NetworkPolicy
metadata:
  name: allow-to-purple-on-8080
spec:
  podSelector:
    matchLabels:
      color: purple
  ingress:
    - ports:
        - protocol: tcp
          port: 8080
```

services provide internal load-balancing and service discovery across pods

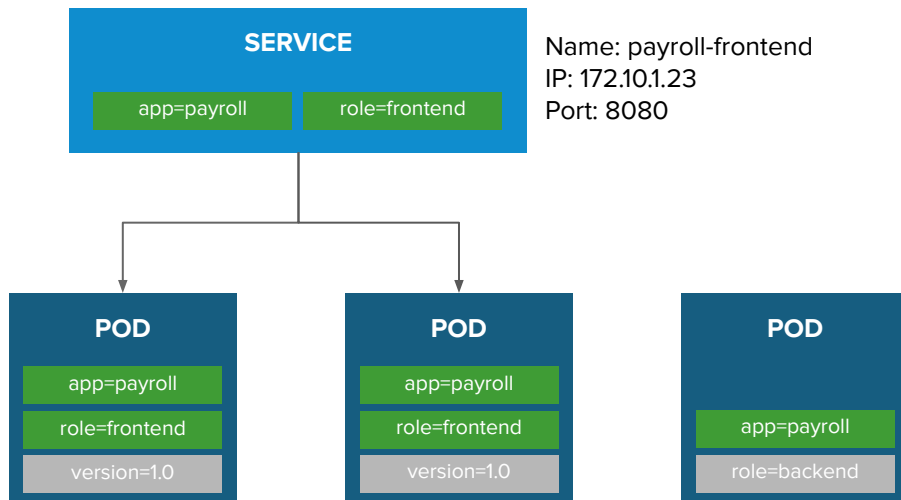


apps can talk to each other via services



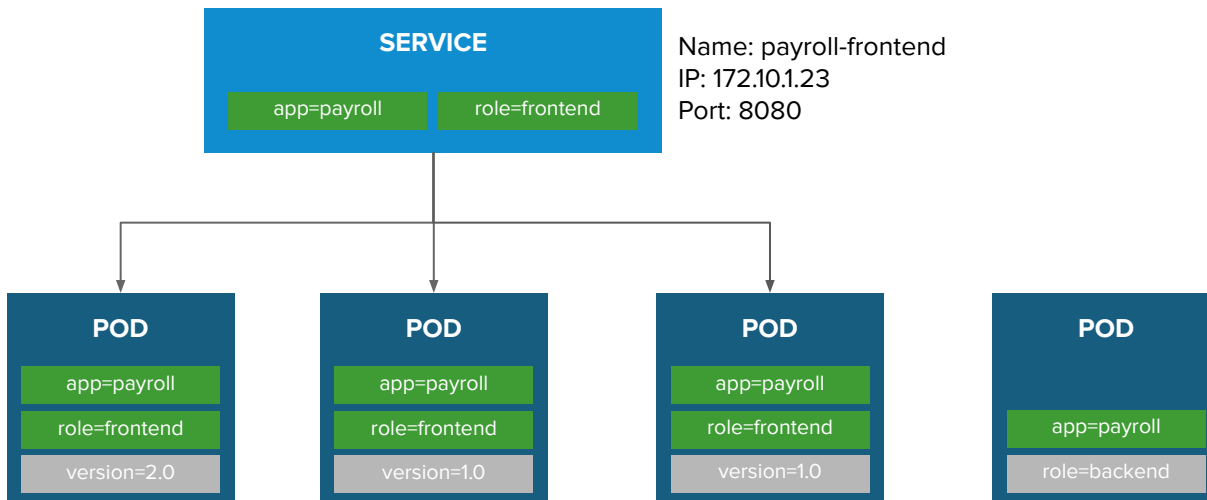
BUILT-IN SERVICE DISCOVERY

INTERNAL LOAD-BALANCING




BUILT-IN SERVICE DISCOVERY

INTERNAL LOAD-BALANCING



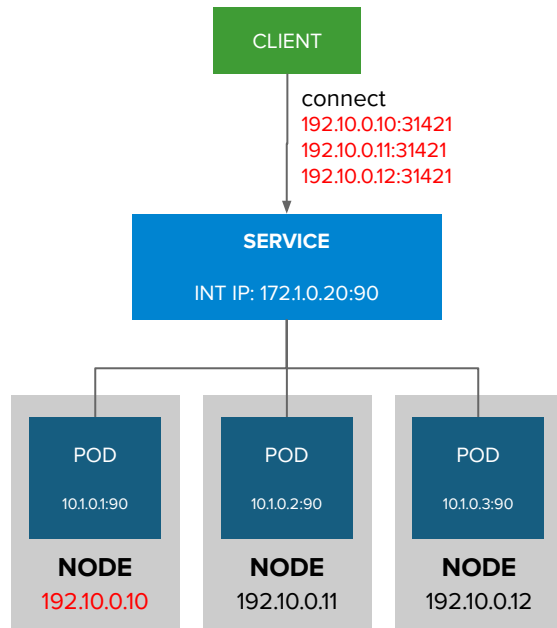
ROUTE SPLIT TRAFFIC

Split Traffic Between
Multiple Services For A/B
Testing, Blue/Green and
Canary Deployments




EXTERNAL TRAFFIC TO A SERVICE ON A RANDOM PORT WITH NODEPORT

- NodePort binds a service to a unique port on all the nodes
- Traffic received on any node redirects to a node with the running service
- Ports in 30K-60K range which usually differs from the service
- Firewall rules must allow traffic to all nodes on the specific port

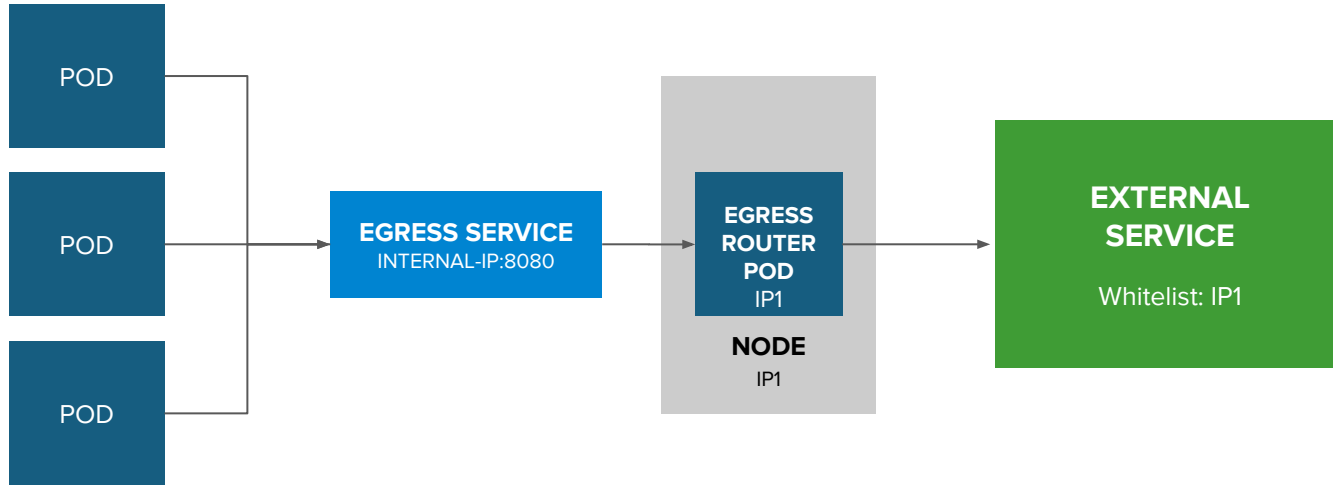


EXTERNAL TRAFFIC TO A SERVICE ON ANY PORT WITH INGRESS

- Access a service with an external IP on any TCP/UDP port, such as
 - Databases
 - Message Brokers
- Automatic IP allocation from a predefined pool using Ingress IP Self-Service
- IP failover pods provide high availability for the IP pool




CONTROL OUTGOING TRAFFIC SOURCE IP WITH EGRESS ROUTER

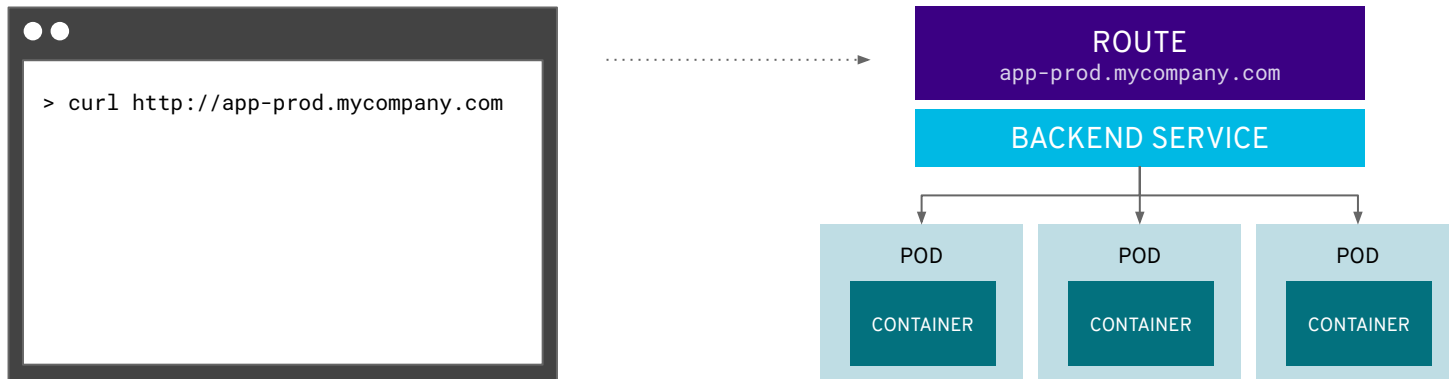


ROUTING AND EXTERNAL LOAD-BALANCING


- Pluggable routing architecture
 - HAProxy Router
 - F5 Router
- Multiple-routers with traffic sharding
- Router supported protocols
 - HTTP/HTTPS
 - WebSockets
 - TLS with SNI
- Non-standard ports via cloud load-balancers, external IP, and NodePort



routes add services to the external load-balancer and provide readable urls for the app

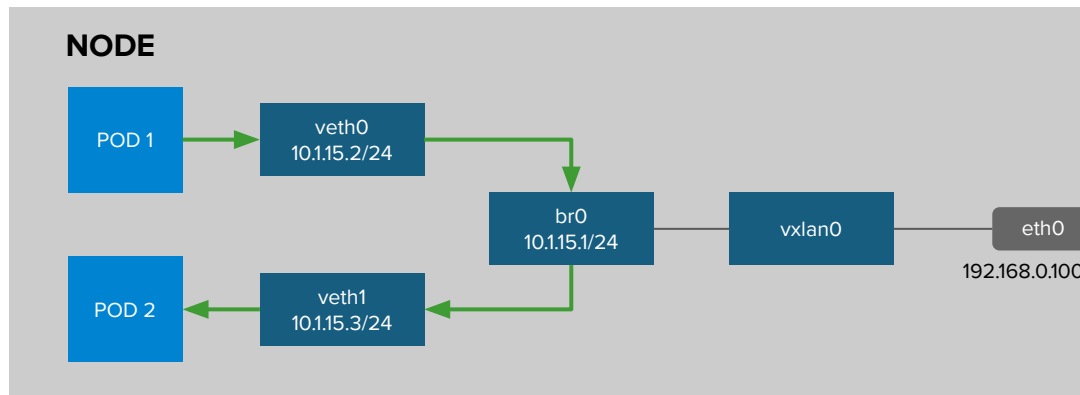


ROUTE EXPOSES SERVICES EXTERNALLY




OPENSSHIFT SDN - OVS PACKET FLOW

Container to Container on the Same Host




OPENSIFT SDN - OVS PACKET FLOW

Container to Container on the Different Hosts



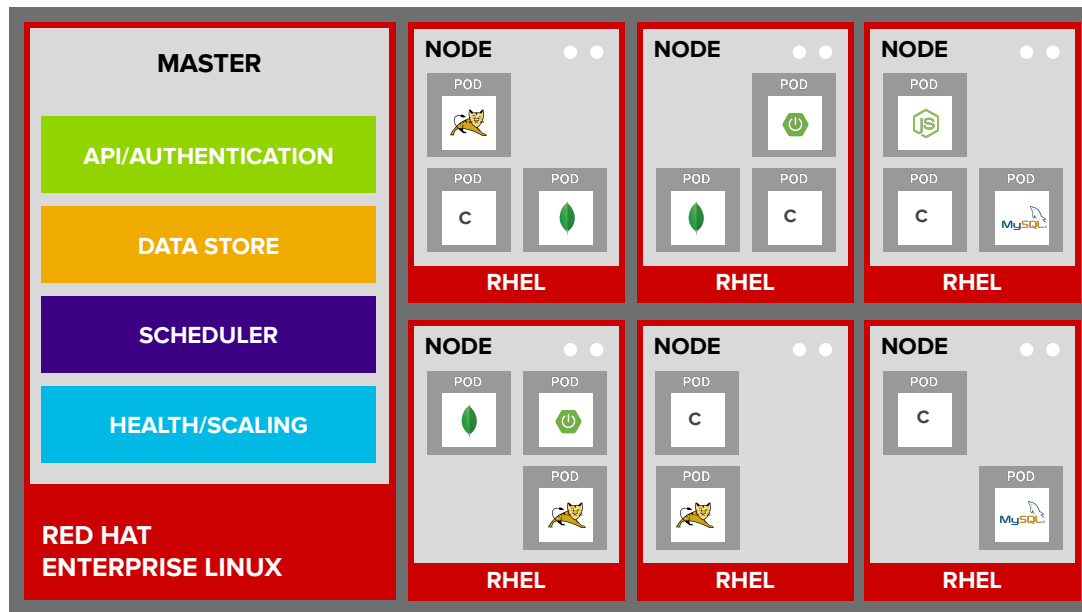
OPENSSHIFT SDN - OVS PACKET FLOW

Container Connects to External Host

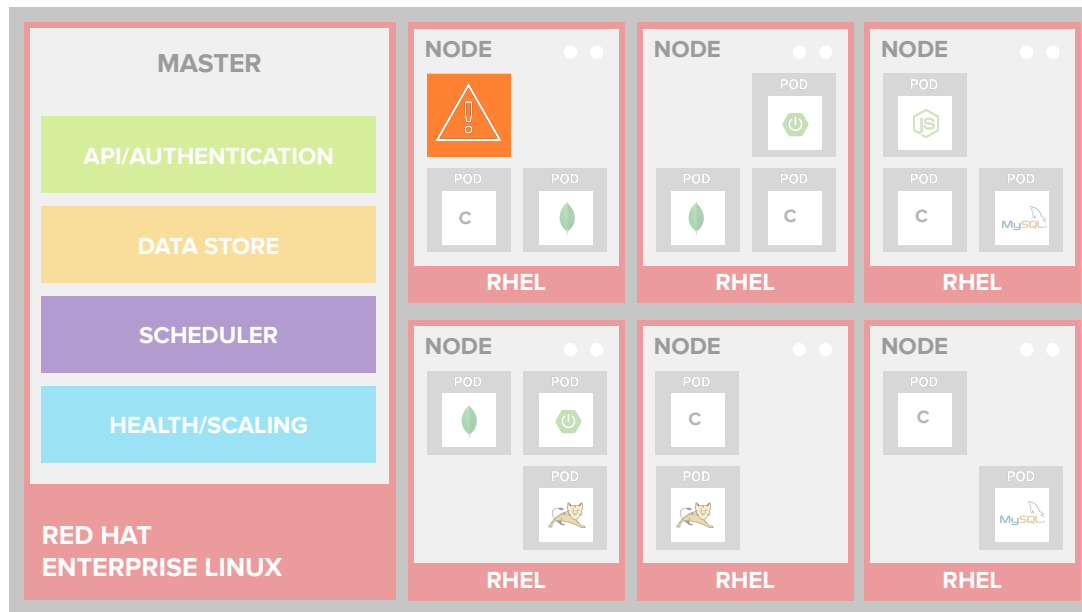


OpenShift Monitoring / Clustering


AUTO-HEALING FAILED PODS



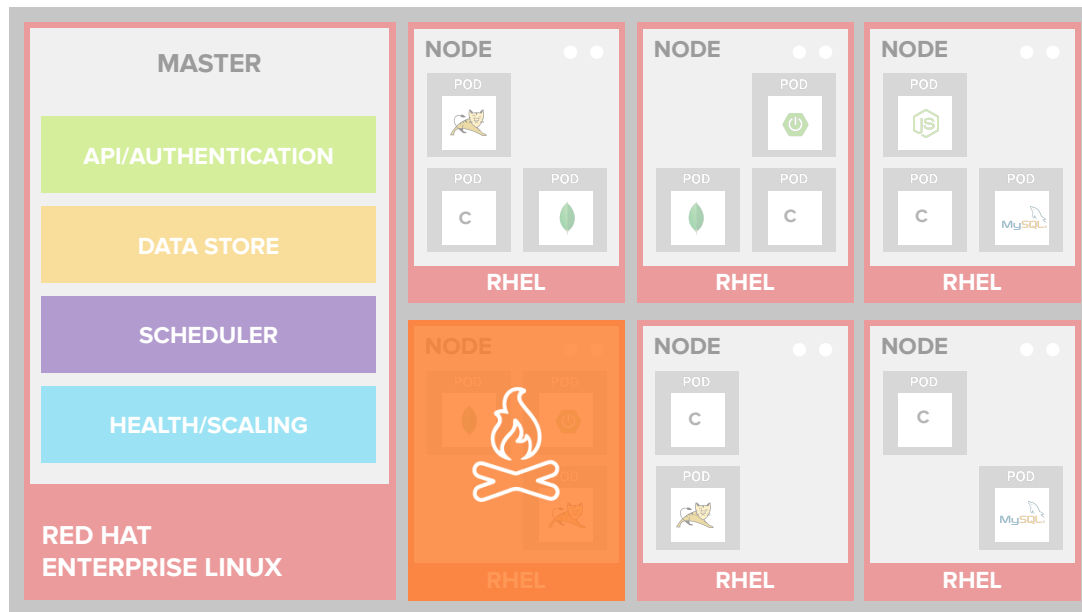
AUTO-HEALING FAILED CONTAINERS



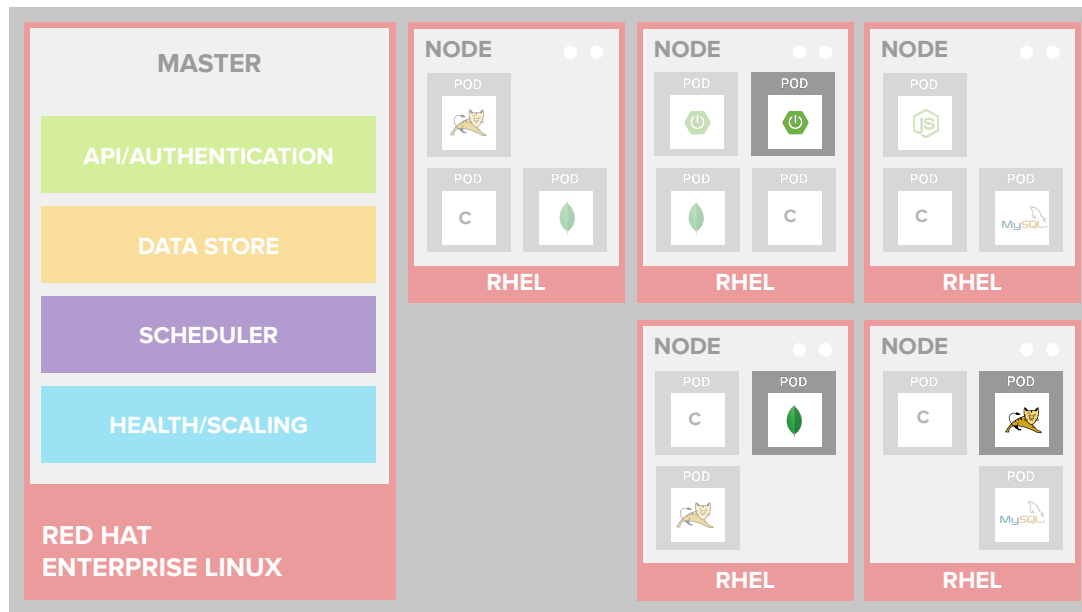
AUTO-HEALING FAILED CONTAINERS



AUTO-HEALING FAILED CONTAINERS



AUTO-HEALING FAILED CONTAINERS



OpenShift persistent Storage

PERSISTENT STORAGE

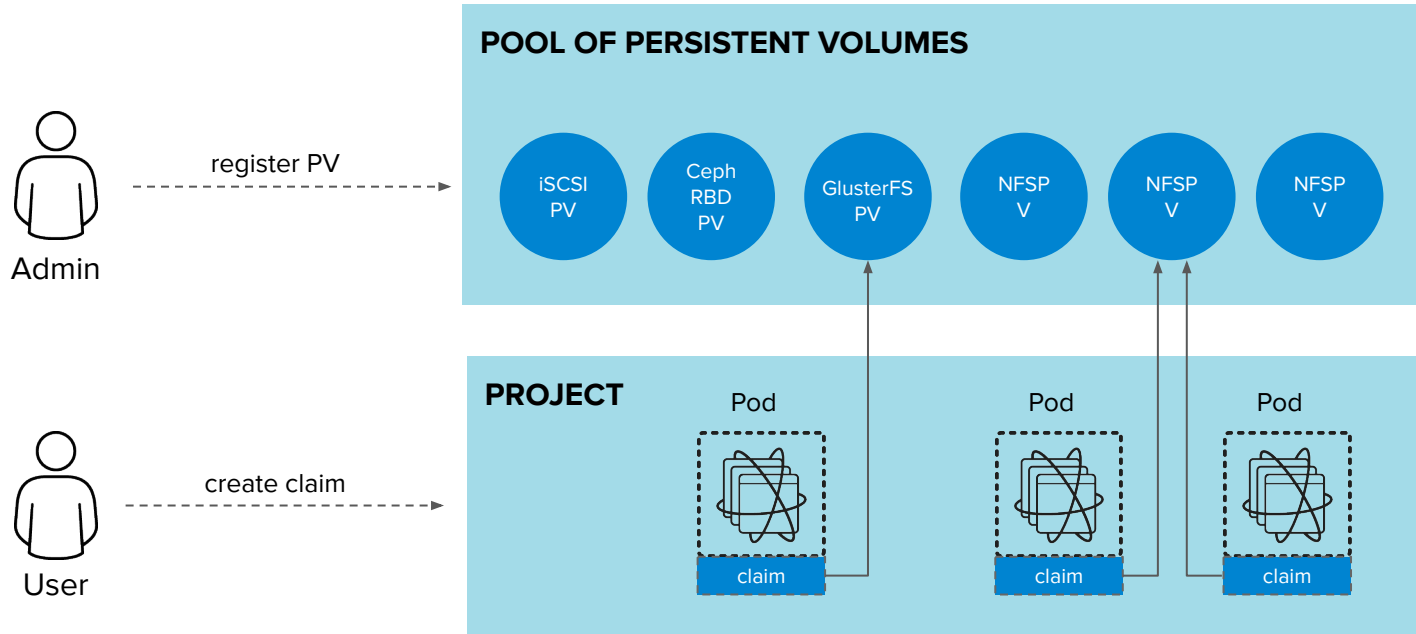
- Persistent Volume (PV) is tied to a piece of network storage
- Provisioned by an administrator (static or dynamically)
- Allows admins to describe storage and users to request storage
- Assigned to pods based on the requested size, access mode, labels and type

NFS	OpenStack Cinder	iSCSI	Azure Disk	AWS EBS	FlexVolume
GlusterFS	Ceph RBD	Fiber Channel	Azure File	GCE Persistent Disk	VMWare vSphere VMDK
		NetApp Trident*	Container Storage Interface (CSI)**		

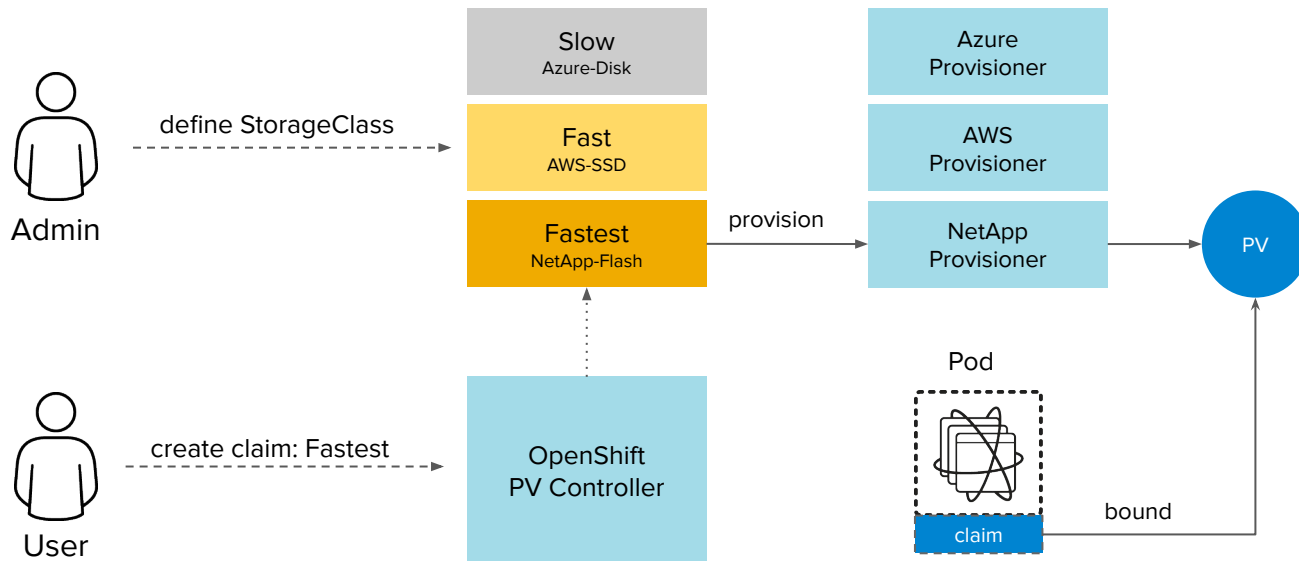
* Shipped and supported by NetApp via TSANet

** Tech Preview

PERSISTENT STORAGE



DYNAMIC VOLUME PROVISIONING



Thank you !